

# 1.8 AWS Database Services

## AWS Cloud Database Offerings

- ❖ AWS Cloud
  - Databases
    - Amazon Elastic Compute Cloud (EC2)
    - AWS Database Migration Service (DMS)
- ❖ AWS Databases
  - Data Warehouse databases
  - Graph databases
  - In-memory data store databases
  - Ledger databases: cryptographically verifiable history of transactions
  - Managed relational databases
  - Non-relational databases
  - Time-series databases

Product	Application	Type
Relational	Transactional	Amazon Aurora
		Amazon RDS
NoSQL database	Internet-scale	Amazon DynamoDB
Non-relational database		Amazon DocumentDB
Data warehouse	Analytic	Amazon Redshift
In-memory data store	Real-time	Amazon ElastiCache
Graph database	Connected data (news feeds)	Amazon Neptune
Time series database	Data collection (IoT devices)	Amazon Timestream
Ledger database	Record of transactions	Amazon QLDB

## **Relational Databases**

- Collection of data items
- Predefined relationships
- Organized as sets of tables
- Store information about objects

### Relational Database Characteristics

- Structured query language (SQL)
- Data integrity
- Transactions: atomic, executed as one single unit, commit or rollback
- Atomic, consistent, isolated, and durable (ACID)

### Relational Database

- Managed: Amazon RDS
- Unmanaged

## **Amazon Relational Database Service (RDS)**

- Procurement
- Configuration
- Backup
- Security
- Availability

### Amazon RDS: Relational database engines

- Amazon Aurora
- Oracle
- Microsoft SQL Server
- PostgreSQL
- MySQL
- MariaDB

### Amazon RDS features

- Software patching

- Vertical scaling: for right intensive applications
- Storage scaling: 3 types – General purpose SSD, provisioned IOPS and magnetic storage
- Horizontal scaling: read replicas, asynchronous

RDS, MySQL, PostgreSQL – 5 read replicas

Amazon Aurora – 15 read replicas

RDS oracle and SQL server – do not support

- Backups: point-to-time, snapshots... 7 to 35 days
- Snapshots
- Multi-AZ deployments
- Encryption: SSL
- IAM authentication
- Amazon CloudWatch: to notify

## **Nonrelational Databases**

### NoSQL Database

- Optimized for scalable performance
- Schema-less data models
- Ease of deployment
- Low latency
- Resilience

### Reviewing SQL versus NoSQL

- Databases reveals that for data storage, SQL uses rows and columns. NoSQL uses key value pairs, documents, wide column and graph objects.
- For schemas SQL is fixed. NoSQL is dynamic.
- For query we use Structured Query Language with SQL databases and with NoSQL databases, we are more focused on a collection of documents.
- With respect to scalability, SQL databases scale very well vertically, while in NoSQL they scale very well horizontally.
- SQL supports transactions. With NoSQL that support can vary. So, it's very flexible.

- For consistency, with the SQL database, we have strong consistency, but with NoSQL database, we can have eventual and strong consistency. So again, quite flexible.

### Using NoSQL Databases

- Big data
- Mobile
- Web applications

### Types of NoSQL databases

- Columnar: read and write columns of data
- Documents: semi-structured, JSON and XML
- Graph: vertices and edges
- In-memory key-value: very low latency, used for heavy workloads

## **Amazon DynamoDB**

- Fully managed cloud database
- Fast and flexible
- Reliable performance
- Automatic scaling
- Consistent
- Can use global tables

### Components of DynamoDB

- Tables
- Items
- Attributes
- Primary key

### DynamoDB Secondary Index

- Primary key attributes
- Secondary key attributes
- Subset of other attributes (optional)
- Can define 5 local and global indexes per table

## In-memory Data Stores

- Used for caching
- Real-time workloads
- In-memory key-value database options
  - ~ Amazon EC2
  - ~ Amazon EBS
  - ~ Amazon ElastiCache

### Caching

- Improved performance
- Data retrieval
- Reduced loads
- Eliminates hotspots
- High throughput
- Low-latency access
- Data is stored in memory

### Strategies of Caching

- Cache hit: when info is requested, contains info requested
- Cache miss: does not contain requested information
- Lazy loading: loading data into cache when necessary
- Write through: data is added to cache when written

### Amazon In-memory Key-Value Data Stores

- In-memory stores
  - Efficient databases
  - Built to be scalable
- Distributed cache
  - More cost-effective
  - Faster performance

### Amazon ElastiCache

- Redis
- Memcached

### Amazon DynamoDB Accelerator

- Fully managed

- Highly available
- In-memory acceleration

## Cloud Database Migration

- Used to migrate databases quickly and securely
- Source database remains fully operational
- Minimizes downtime

### Using AWS DMS to perform Migration

- Loading existing data
- Applying any cached changes
- Performing ongoing replication

### AWS Data Migration

- Homogenous:
- Heterogeneous: use AWS Schema Conversion Tool (SCT)
  - automatically converts database schema
  - scans application source code
  - performs cloud-native code optimization

## Installing the Python SDK and PyMySQL

Download latest version of python

Then install boto3

```

PROBLEMS  OUTPUT  DEBUG CONSOLE  TERMINAL
1: pip

PS C:\Users\Joe\aws projects> pip install boto3
Collecting boto3
  Using cached boto3-1.13.21-py2.py3-none-any.whl (128 kB)
Requirement already satisfied: s3transfer<0.4.0,>=0.3.0 in c:\users\joe\appdata\roaming\python\python37\site-packages (from boto3) (0.3.3)
Requirement already satisfied: jmespath<1.0.0,>=0.7.1 in c:\users\joe\appdata\roaming\python\python37\site-packages (from boto3) (0.9.5)
Requirement already satisfied: botocore<1.17.0,>=1.16.21 in c:\users\joe\appdata\local\programs\python\python37-32\lib\site-packages (from boto3) (1.16.21)
Requirement already satisfied: docutils<0.16,>=0.10 in c:\users\joe\appdata\roaming\python\python37\site-packages (from botocore<1.17.0,>=1.16.21->boto3) (0.15.2)
Requirement already satisfied: urllib3<1.26,>=1.20; python_version != "3.4" in c:\users\joe\appdata\roaming\python\python37\site-packages (from botocore<1.17.0,>=1.16.21->boto3) (1.25.9)
Requirement already satisfied: python-dateutil<3.0.0,>=2.1 in c:\users\joe\appdata\roaming\python\python37\site-packages (from botocore<1.17.0,>=1.16.21->boto3) (2.8.1)
Requirement already satisfied: six>=1.5 in c:\users\joe\appdata\roaming\python\python37\site-packages (from python-dateutil<3.0.0,>=2.1->botocore<1.17.0,>=1.16.21->boto3) (1.14.0)
Installing collected packages: boto3
Successfully installed boto3-1.13.21

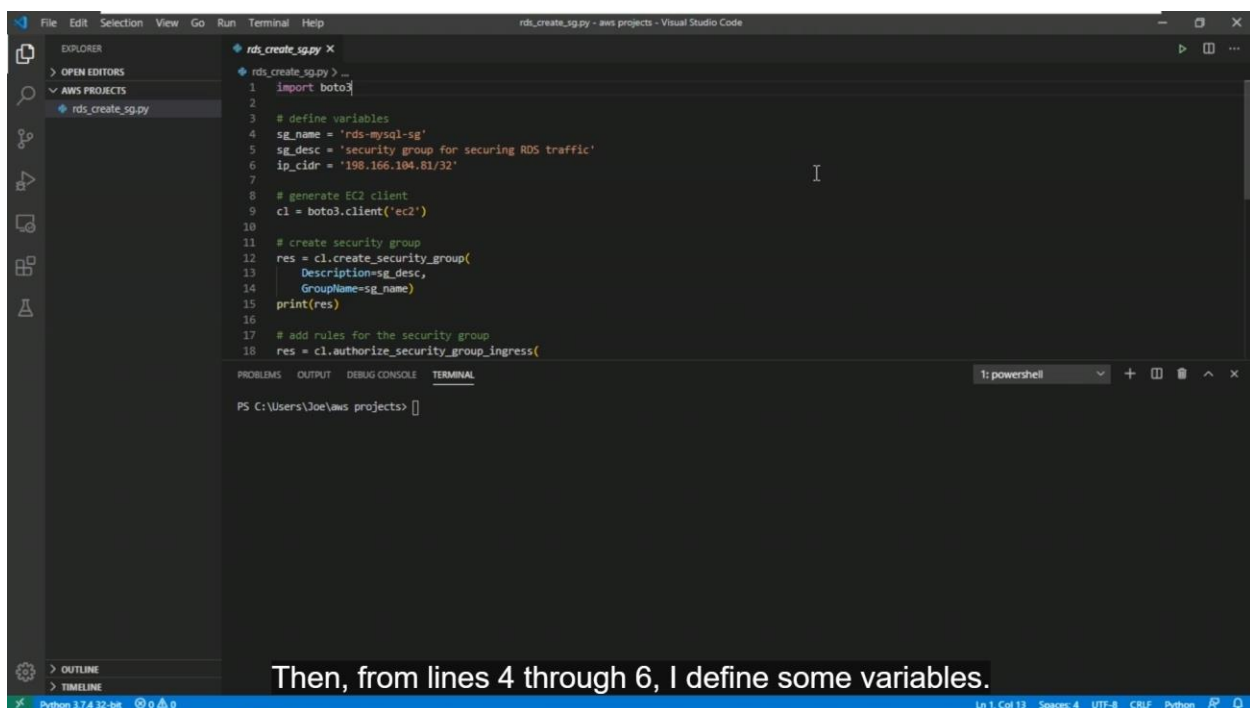
```

```
PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL
1: powershell +

PS C:\Users\Joe\aws projects> pip install pymysql
Collecting pymysql
  Using cached PyMySQL-0.9.3-py2.py3-none-any.whl (47 kB)
Installing collected packages: pymysql
Successfully installed pymysql-0.9.3
PS C:\Users\Joe\aws projects> |
```

pip= python packages

## Generating Security Groups for Amazon RDS



```
File Edit Selection View Go Run Terminal Help
rds_create_sg.py - aws projects - Visual Studio Code

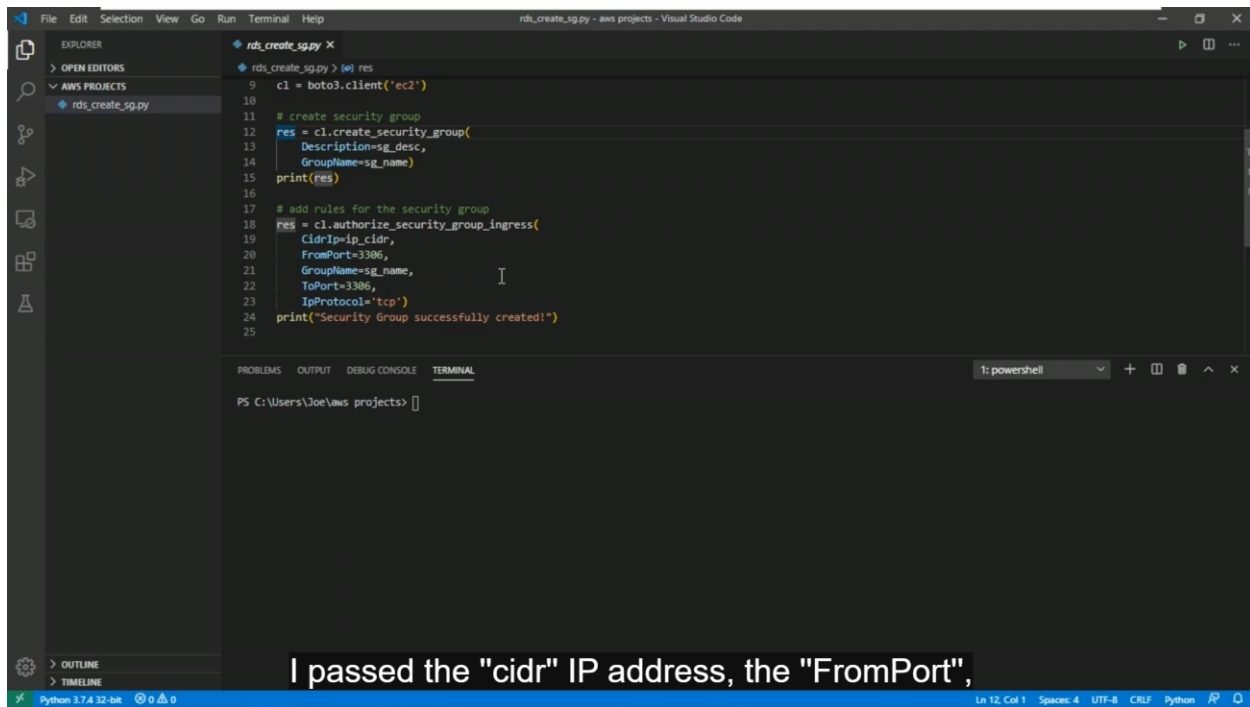
EXPLORER
  OPEN EDITORS
  AWS PROJECTS
    rds_create_sg.py

rds_create_sg.py X
1 import boto3
2
3 # define variables
4 sg_name = 'rds-mysql-sg'
5 sg_desc = 'security group for securing RDS traffic'
6 ip_cidr = '198.166.104.81/32'
7
8 # generate EC2 client
9 cl = boto3.client('ec2')
10
11 # create security group
12 res = cl.create_security_group(
13     Description=sg_desc,
14     GroupName=sg_name)
15 print(res)
16
17 # add rules for the security group
18 res = cl.authorize_security_group_ingress(

PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL
1: powershell +

PS C:\Users\Joe\aws projects> |
```

Then, from lines 4 through 6, I define some variables.

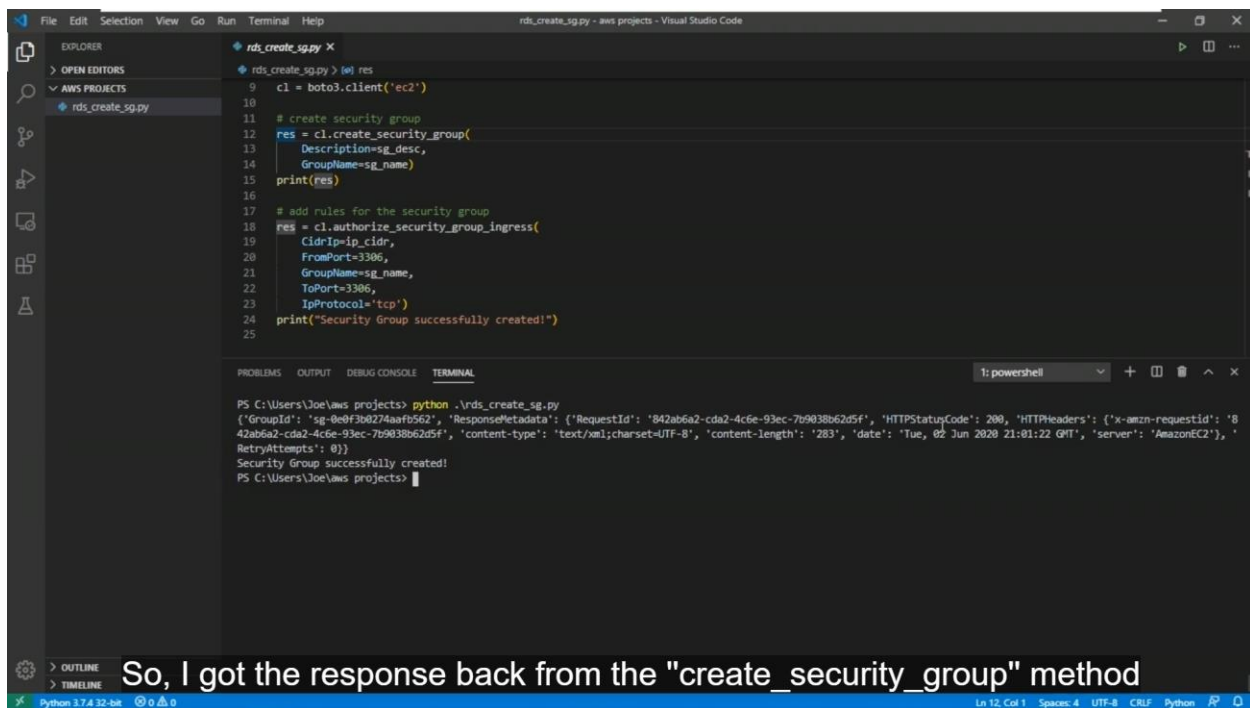


The screenshot shows the Visual Studio Code interface with a file named `rds_create_sg.py` open in the editor. The script contains the following Python code:

```
9  c1 = boto3.client('ec2')
10
11  # create security group
12  res = c1.create_security_group(
13      Description=sg_desc,
14      GroupName=sg_name)
15  print(res)
16
17  # add rules for the security group
18  res = c1.authorize_security_group_ingress(
19      CidrIp=ip_cidr,
20      FromPort=3306,
21      GroupName=sg_name,
22      ToPort=3306,
23      IpProtocol='tcp')
24  print("Security Group successfully created!")
25
```

The terminal window at the bottom shows the command prompt `PS C:\Users\joe\aws projects>` with a cursor.

I passed the "cidr" IP address, the "FromPort",



The screenshot shows the same Visual Studio Code interface, but the terminal window now displays the output of running the script:

```
PS C:\Users\joe\aws projects> python .\rds_create_sg.py
{'GroupId': 'sg-0e0f3b0274aafb562', 'ResponseMetadata': {'RequestId': '842ab6a2-cda2-4c6e-93ec-7b9038b62d5f', 'HTTPStatusCode': 200, 'HTTPHeaders': {'x-amzn-requestid': '842ab6a2-cda2-4c6e-93ec-7b9038b62d5f', 'content-type': 'text/xml; charset=UTF-8', 'content-length': '283', 'date': 'Tue, 02 Jun 2020 21:01:22 GMT', 'server': 'AmazonEC2'}, 'RetryAttempts': 0}}
Security Group successfully created!
PS C:\Users\joe\aws projects>
```

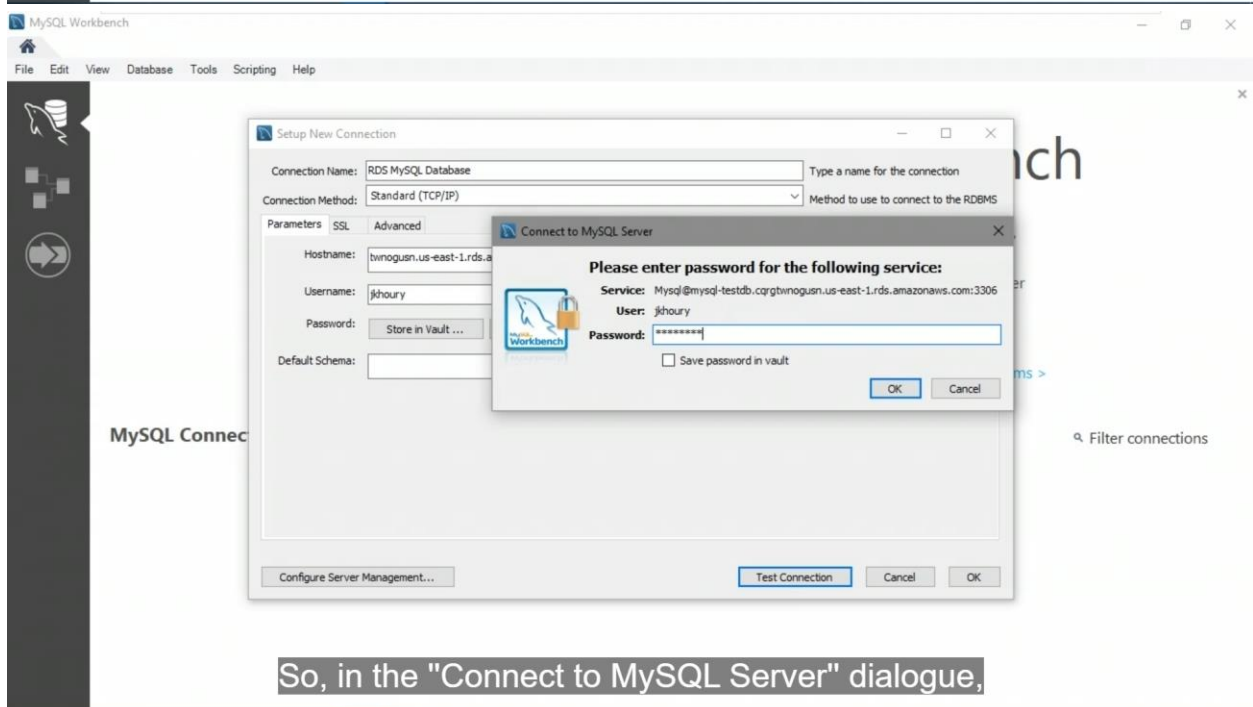
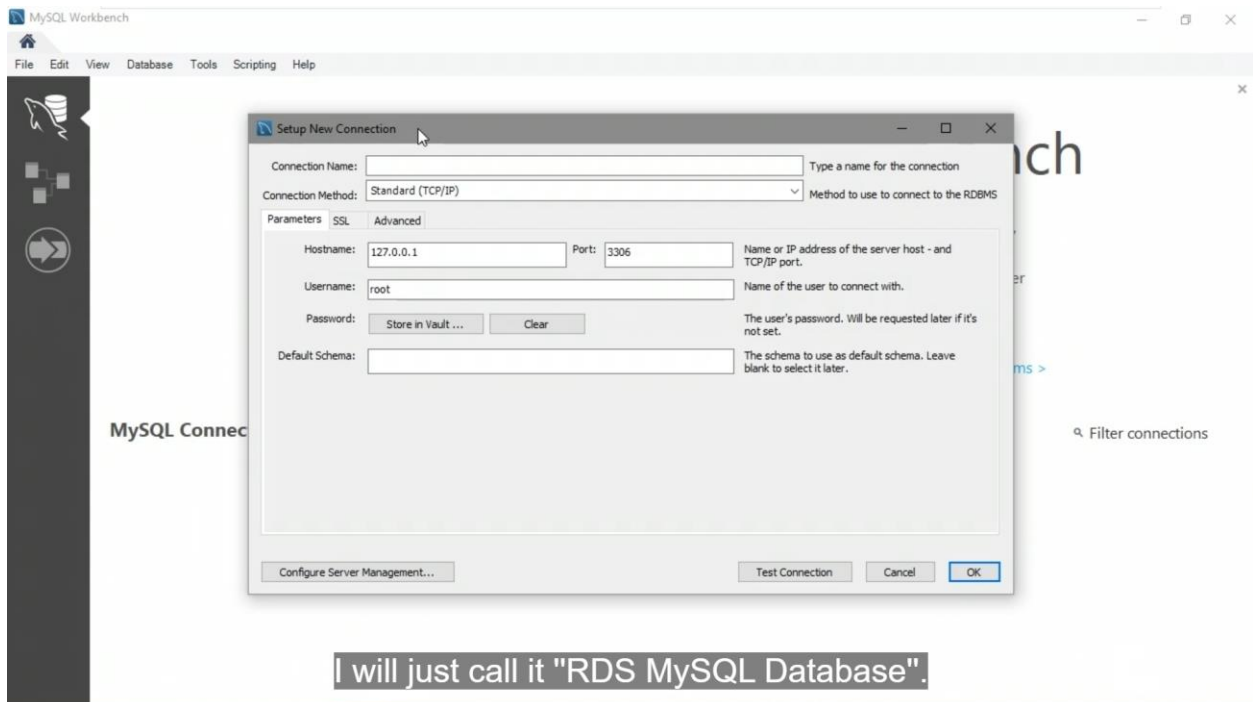
So, I got the response back from the "create\_security\_group" method

## Creating an Amazon RDS Database Instance

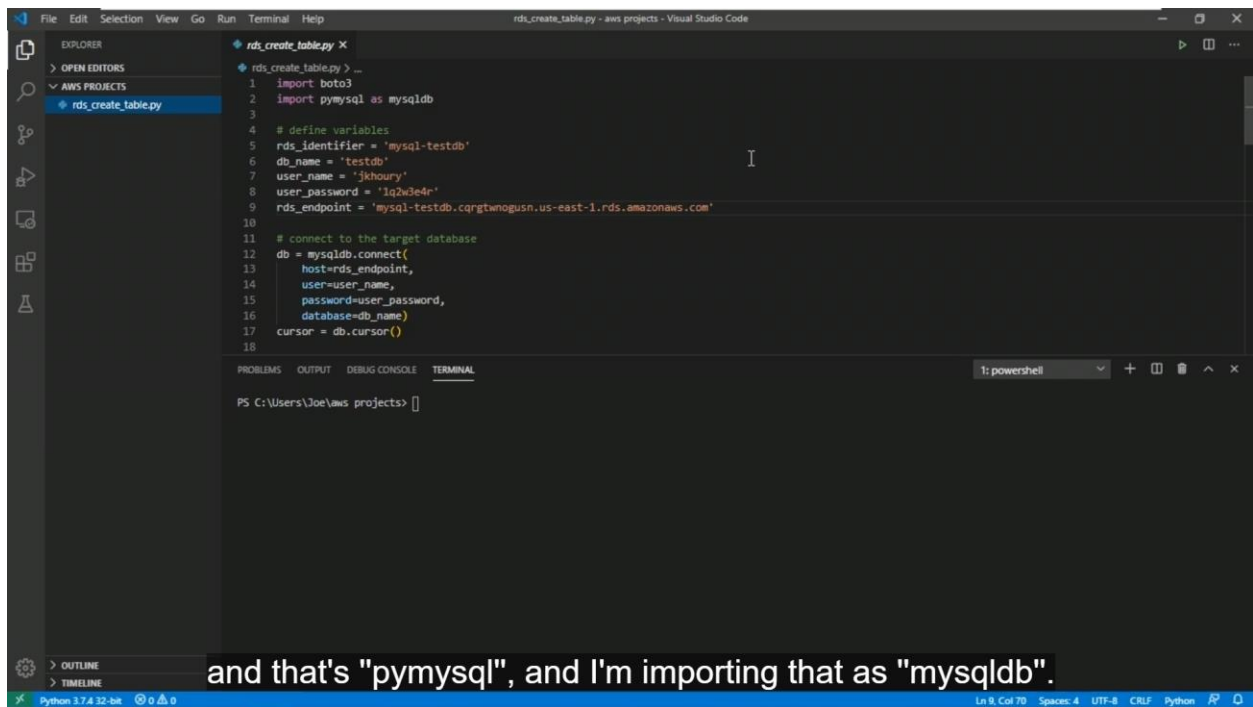


# Connecting to a Database Using MySQL Workbench

Download MySQL workbench



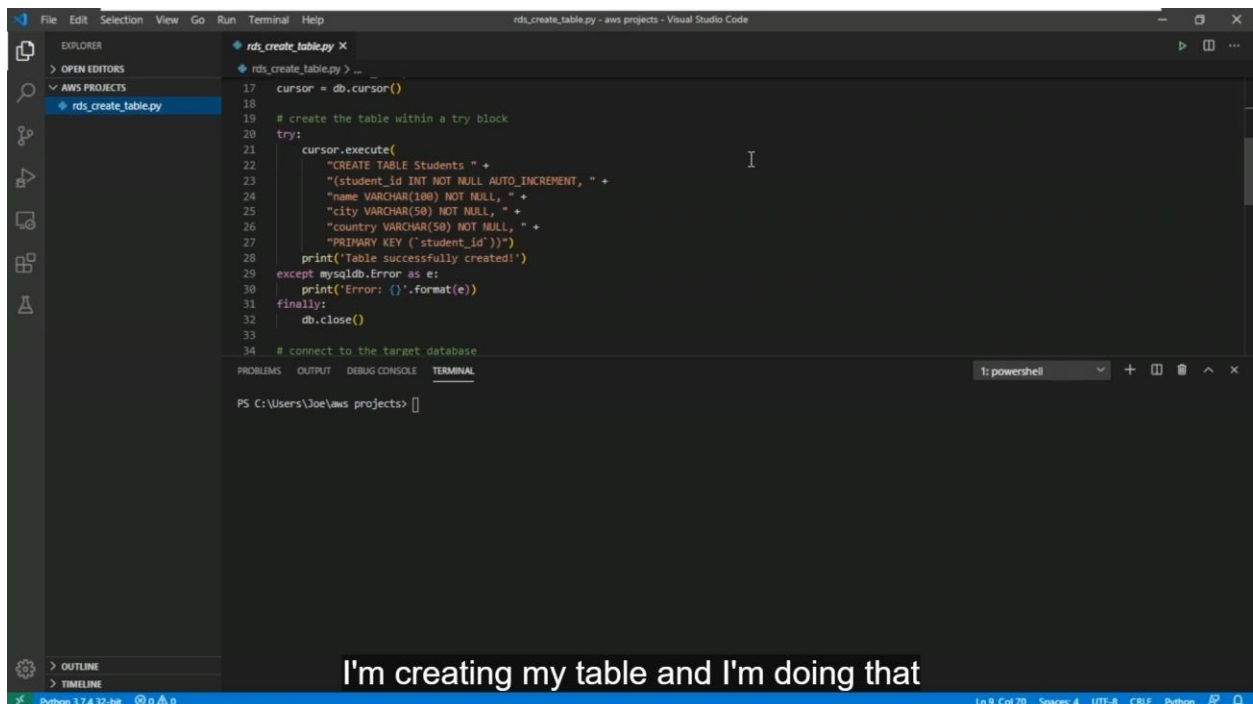
## Creating Tables in an Amazon RDS MySQL Database



The screenshot shows a Visual Studio Code editor with a file named `rds_create_table.py`. The code imports `boto3` and `pymysql` as `mysqldb`. It defines variables for `rds_identifier`, `db_name`, `user_name`, `user_password`, and `rds_endpoint`. It then connects to the database using `mysqldb.connect()` and creates a cursor.

```
1 import boto3
2 import pymysql as mysqldb
3
4 # define variables
5 rds_identifier = 'mysql-testdb'
6 db_name = 'testdb'
7 user_name = 'jkhoury'
8 user_password = 'lq2w3e4r'
9 rds_endpoint = 'mysql-testdb.cqrgtwnogusn.us-east-1.rds.amazonaws.com'
10
11 # connect to the target database
12 db = mysqldb.connect(
13     host=rds_endpoint,
14     user=user_name,
15     password=user_password,
16     database=db_name)
17 cursor = db.cursor()
18
```

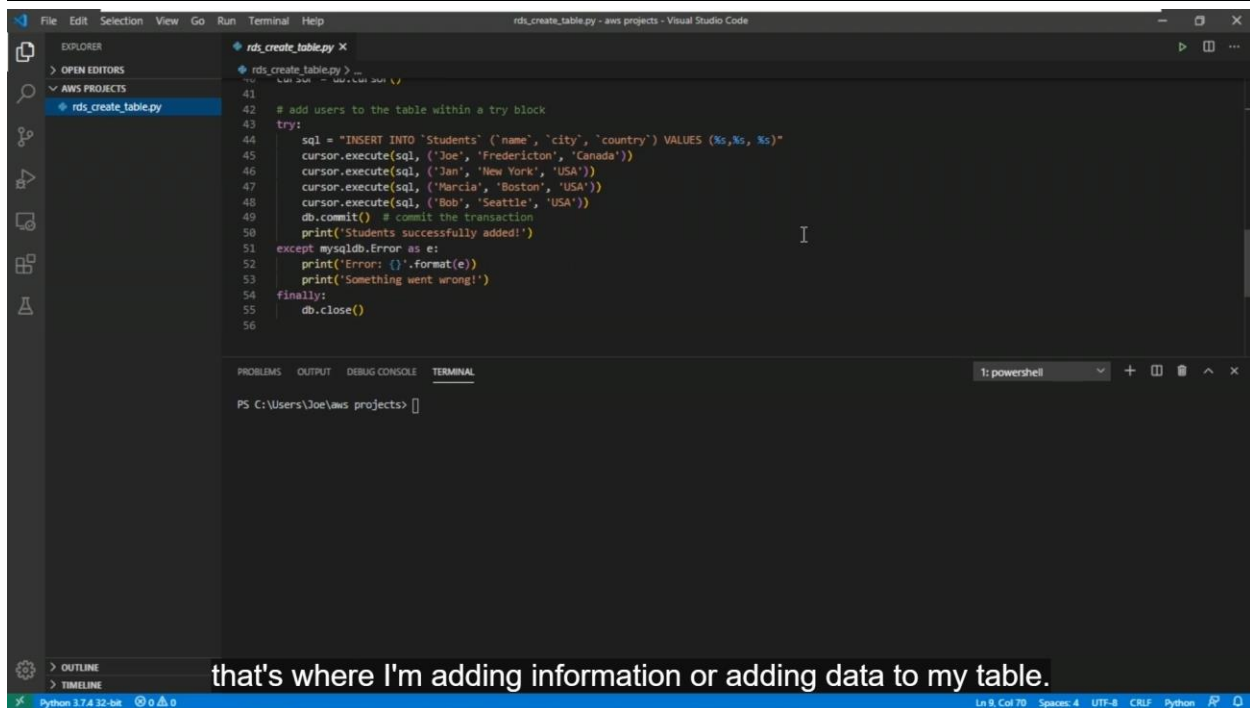
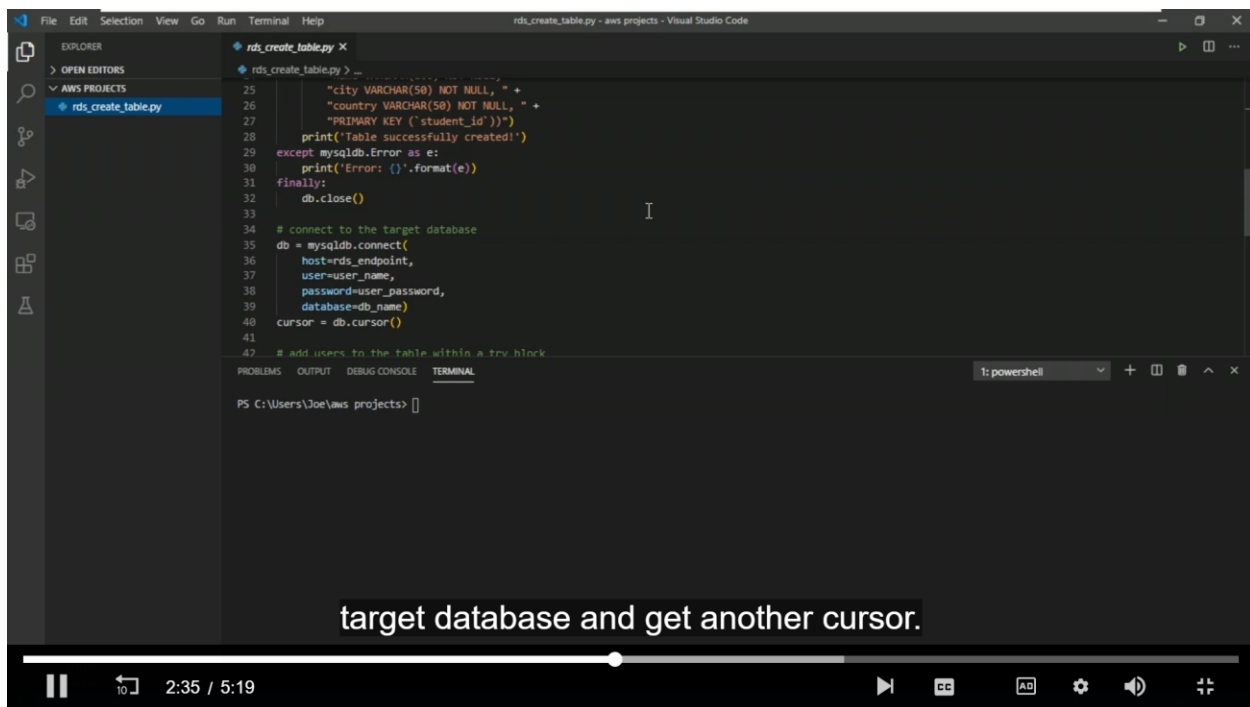
Below the code editor is a terminal window showing the PowerShell prompt: `PS C:\Users\Joe\aws projects>`. A text overlay at the bottom of the screenshot reads: "and that's 'pymysql', and I'm importing that as 'mysqldb'".

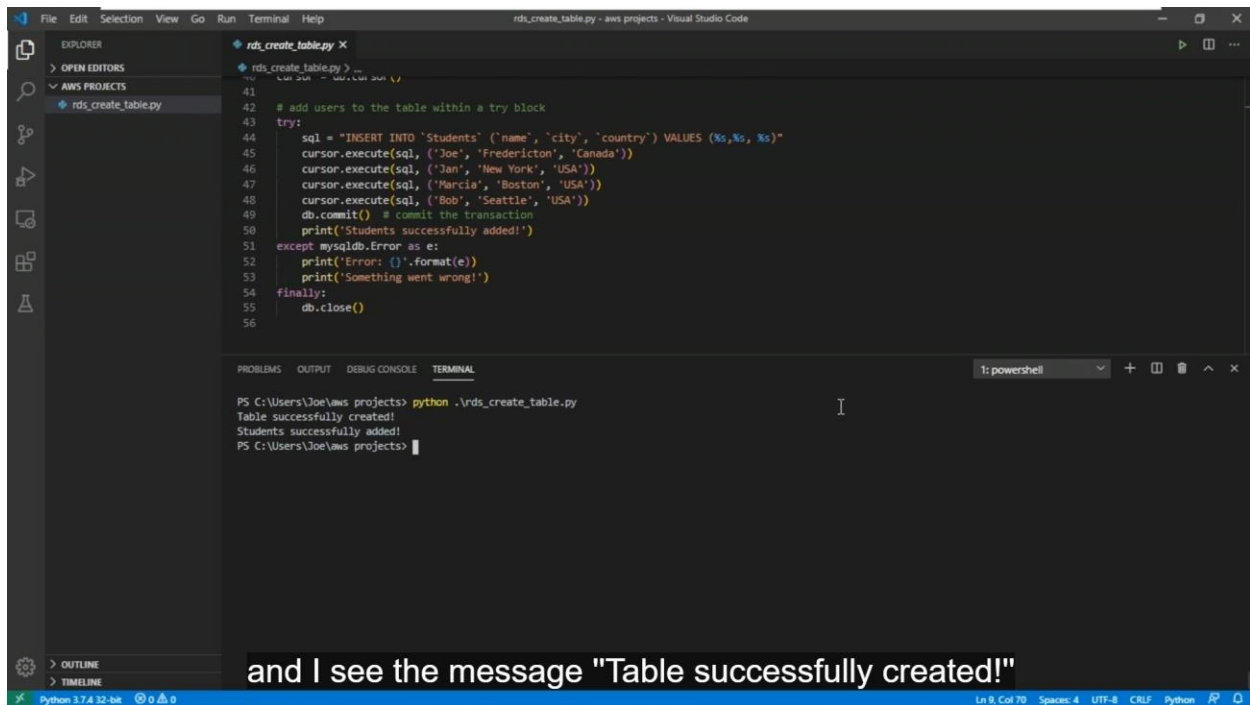


The screenshot shows the continuation of the `rds_create_table.py` script. It uses a `try` block to execute a `CREATE TABLE` SQL statement. The table has columns for `student_id`, `name`, `city`, and `country`. It also includes a `finally` block to close the database connection.

```
17 cursor = db.cursor()
18
19 # create the table within a try block
20 try:
21     cursor.execute(
22         "CREATE TABLE Students " +
23         "(student_id INT NOT NULL AUTO_INCREMENT, " +
24         "name VARCHAR(100) NOT NULL, " +
25         "city VARCHAR(50) NOT NULL, " +
26         "country VARCHAR(50) NOT NULL, " +
27         "PRIMARY KEY ('student_id'))")
28     print('Table successfully created!')
29 except mysqldb.Error as e:
30     print('Error: {}'.format(e))
31 finally:
32     db.close()
33
34 # connect to the target database
```

Below the code editor is a terminal window showing the PowerShell prompt: `PS C:\Users\Joe\aws projects>`. A text overlay at the bottom of the screenshot reads: "I'm creating my table and I'm doing that".





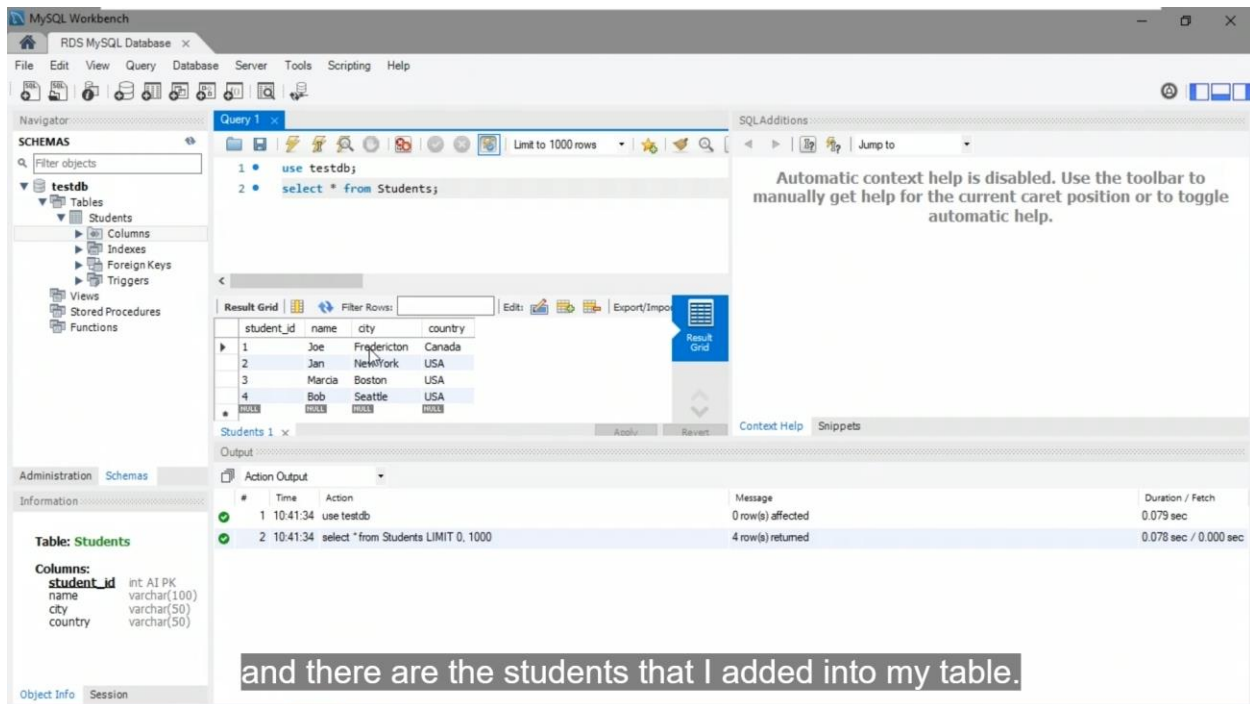
The screenshot shows the Visual Studio Code interface with a file named `rds_create_table.py` open. The script is a Python program that connects to an AWS RDS instance and creates a table named `Students`. It uses a try-except block to handle database errors. The terminal output shows the successful execution of the script, including the message "Table successfully created!".

```
41
42 # add users to the table within a try block
43 try:
44     sql = "INSERT INTO 'Students' ('name', 'city', 'country') VALUES (%s,%s, %s)"
45     cursor.execute(sql, ('Joe', 'Frederickton', 'Canada'))
46     cursor.execute(sql, ('Jan', 'New York', 'USA'))
47     cursor.execute(sql, ('Marcia', 'Boston', 'USA'))
48     cursor.execute(sql, ('Bob', 'Seattle', 'USA'))
49     db.commit() # commit the transaction
50     print('Students successfully added!')
51 except MySQLdb.Error as e:
52     print('Error: {}'.format(e))
53     print('Something went wrong!')
54 finally:
55     db.close()
56
```

PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL

PS C:\Users\Joe\aws projects> python .\rds\_create\_table.py  
Table successfully created!  
Students successfully added!  
PS C:\Users\Joe\aws projects>

and I see the message "Table successfully created!"



The screenshot shows the MySQL Workbench interface. The 'Schemas' pane on the left shows the 'testdb' database selected. The 'Query 1' window shows the SQL query `select * from Students;`. The 'Result Grid' displays the data from the `Students` table. The 'Output' pane shows the execution log, including the message "4 row(s) returned".

student_id	name	city	country
1	Joe	Frederickton	Canada
2	Jan	New York	USA
3	Marcia	Boston	USA
4	Bob	Seattle	USA

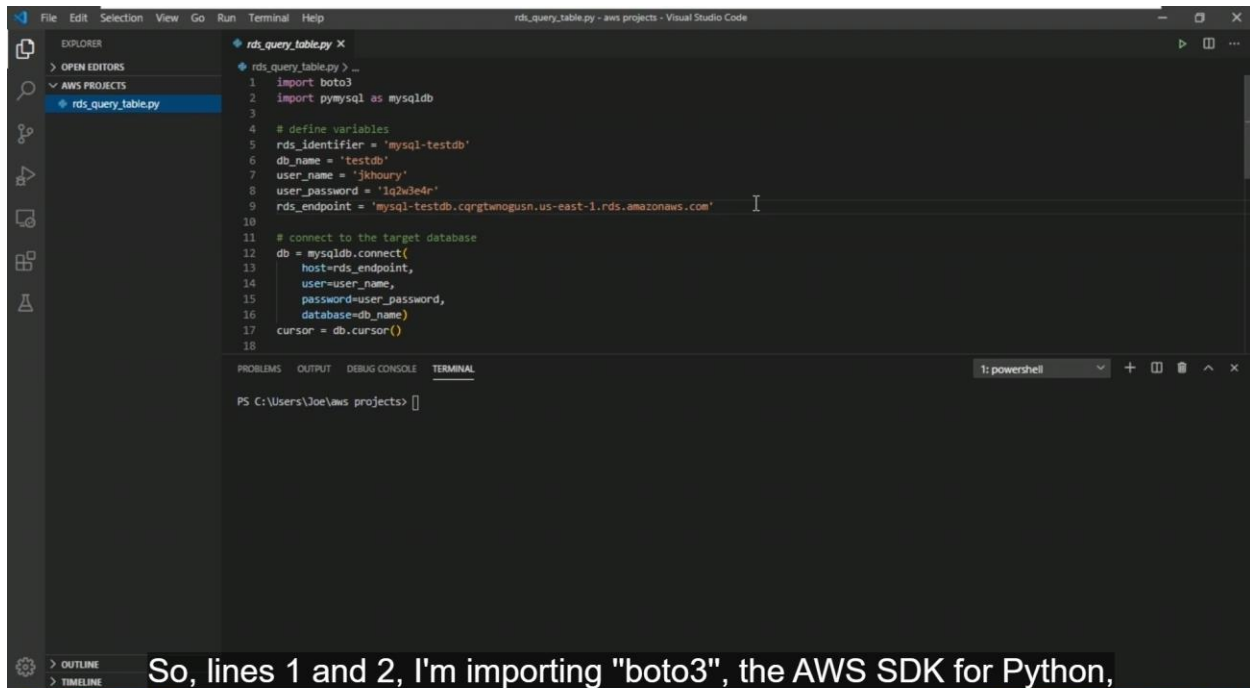
Table: **Students**

Columns:

- `student_id` int AI PK
- `name` varchar(100)
- `city` varchar(50)
- `country` varchar(50)

and there are the students that I added into my table.

# Querying Tables in an Amazon RDS MySQL Database



```
File Edit Selection View Go Run Terminal Help
rds_query_table.py - aws projects - Visual Studio Code

EXPLORER
  > OPEN EDITORS
  > AWS PROJECTS
    rds_query_table.py

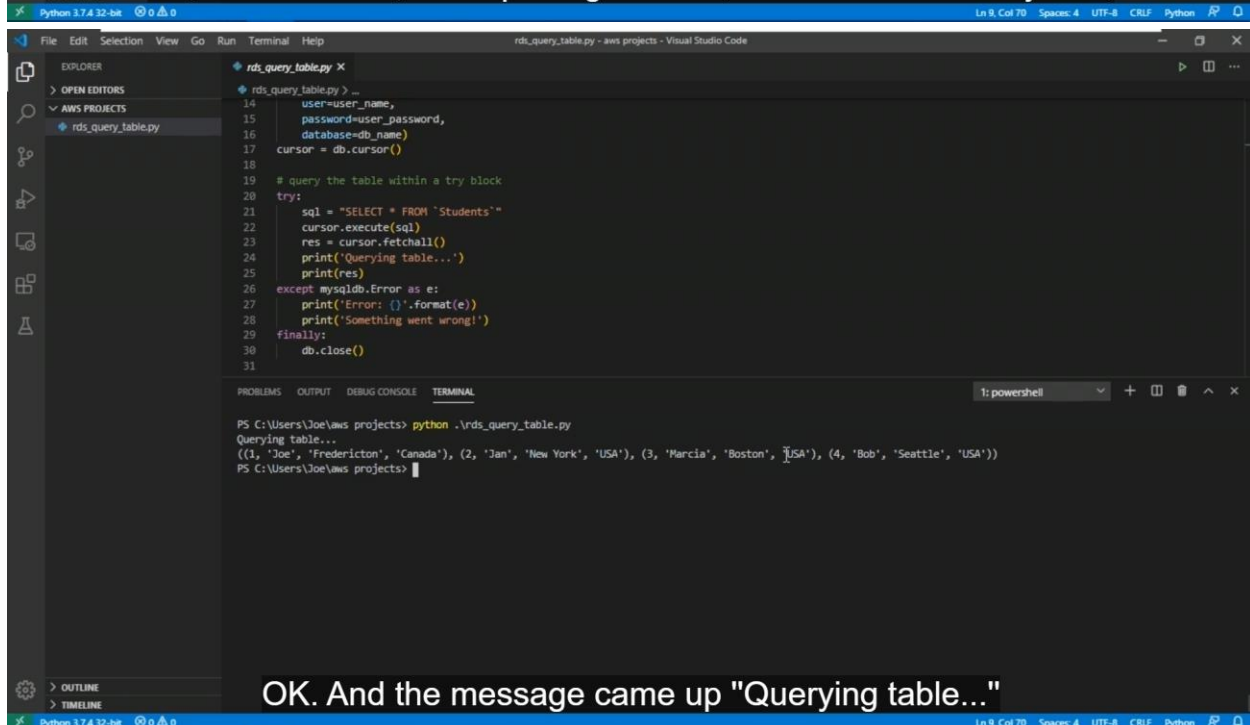
rds_query_table.py X
1 import boto3
2 import pymysql as mysqldb
3
4 # define variables
5 rds_identifier = 'mysql-testdb'
6 db_name = 'testdb'
7 user_name = 'jkhoury'
8 user_password = '1q2w3e4r'
9 rds_endpoint = 'mysql-testdb.cqrgtwmogusn.us-east-1.rds.amazonaws.com'
10
11 # connect to the target database
12 db = mysqldb.connect(
13     host=rds_endpoint,
14     user=user_name,
15     password=user_password,
16     database=db_name)
17 cursor = db.cursor()
18
```

PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL

1: powershell

PS C:\Users\Joe\aws projects>

So, lines 1 and 2, I'm importing "boto3", the AWS SDK for Python,



```
File Edit Selection View Go Run Terminal Help
rds_query_table.py - aws projects - Visual Studio Code

EXPLORER
  > OPEN EDITORS
  > AWS PROJECTS
    rds_query_table.py

rds_query_table.py X
14 user=user_name,
15 password=user_password,
16 database=db_name)
17 cursor = db.cursor()
18
19 # query the table within a try block
20 try:
21     sql = "SELECT * FROM 'Students'"
22     cursor.execute(sql)
23     res = cursor.fetchall()
24     print('Querying table...')
25     print(res)
26 except mysqldb.Error as e:
27     print('Error: {}'.format(e))
28     print('Something went wrong!')
29 finally:
30     db.close()
31
```

PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL

1: powershell

PS C:\Users\Joe\aws projects> python .\rds\_query\_table.py

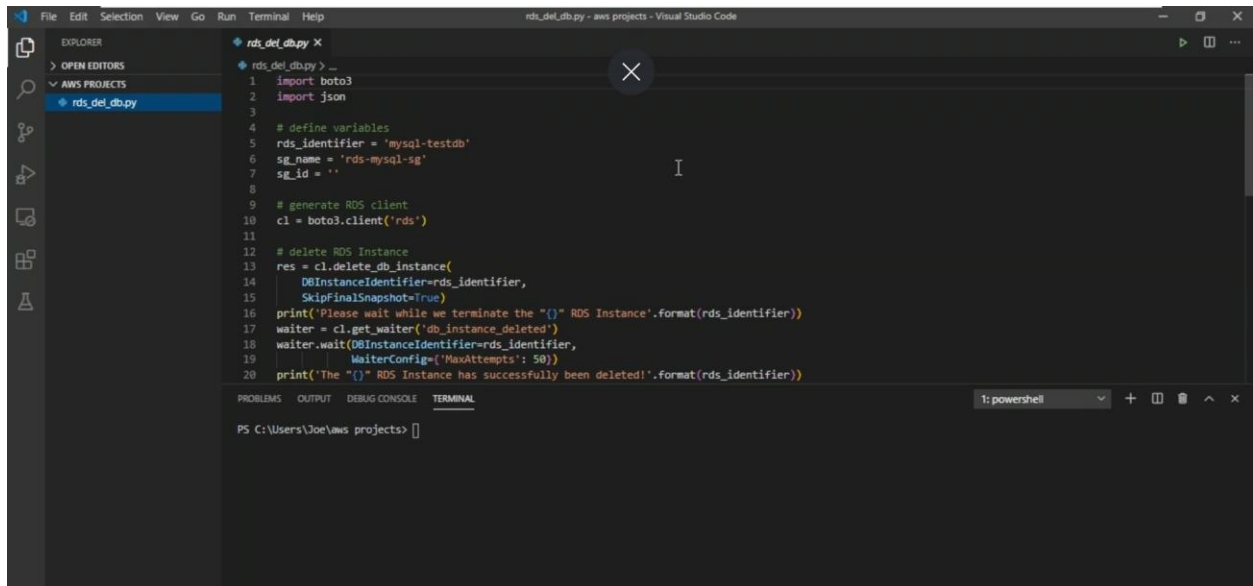
Querying table...

((1, 'Joe', 'Frederickton', 'Canada'), (2, 'Jan', 'New York', 'USA'), (3, 'Marcia', 'Boston', 'USA'), (4, 'Bob', 'Seattle', 'USA'))

PS C:\Users\Joe\aws projects>

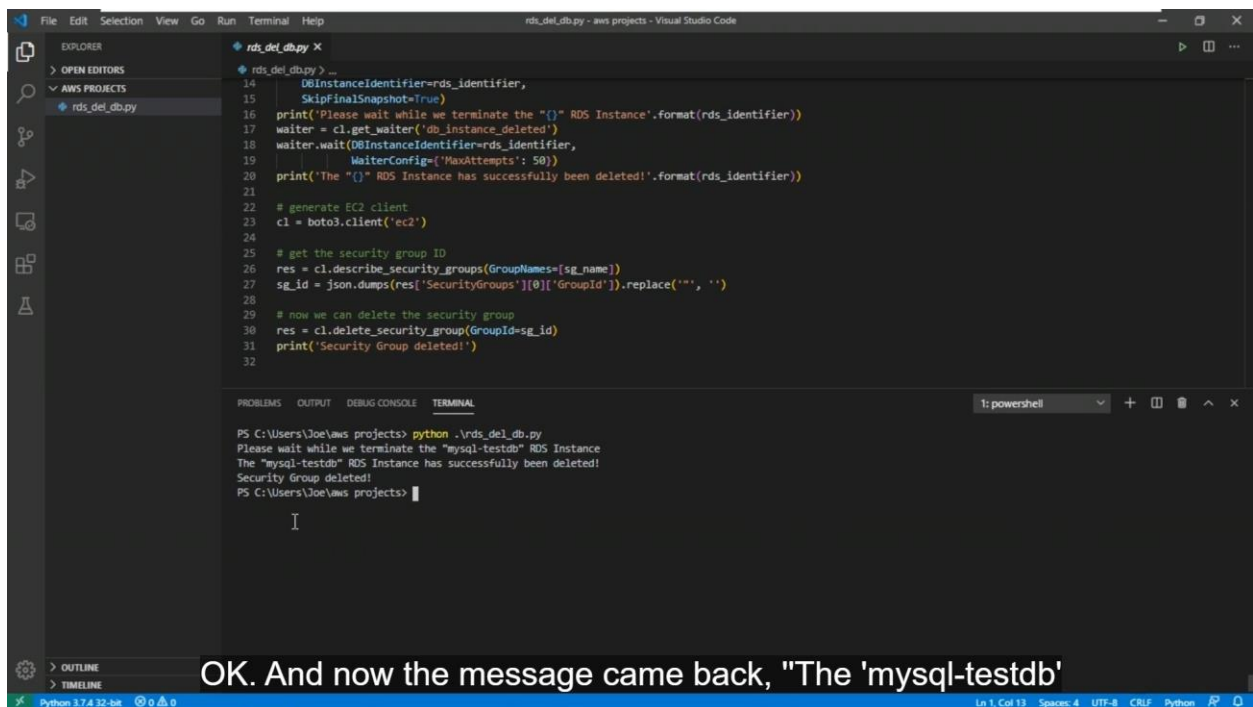
OK. And the message came up "Querying table..."

## Deleting Databases and Security Groups



The screenshot shows the Visual Studio Code editor with the file `rds_del_db.py` open. The script is a Python program that uses the boto3 library to delete an RDS instance. It defines variables for the instance identifier, security group name, and security group ID. It then generates an RDS client and calls `delete_db_instance` with the instance identifier and `SkipFinalSnapshot=True`. The script includes a waiter to ensure the instance is deleted and a final print statement.

```
1 import boto3
2 import json
3
4 # define variables
5 rds_identifier = 'mysql-testdb'
6 sg_name = 'rds-mysql-sg'
7 sg_id = ''
8
9 # generate RDS client
10 cl = boto3.client('rds')
11
12 # delete RDS Instance
13 res = cl.delete_db_instance(
14     DBInstanceIdentifier=rds_identifier,
15     SkipFinalSnapshot=True)
16 print('Please wait while we terminate the "{}" RDS Instance'.format(rds_identifier))
17 waiter = cl.get_waiter('db_instance_deleted')
18 waiter.wait(DBInstanceIdentifier=rds_identifier,
19             WaiterConfig={'MaxAttempts': 50})
20 print('The "{}" RDS Instance has successfully been deleted!'.format(rds_identifier))
```



The screenshot shows the Visual Studio Code editor with the file `rds_del_db.py` open. The script is a Python program that uses the boto3 library to delete an RDS instance and a security group. It defines variables for the instance identifier, security group name, and security group ID. It then generates an RDS client and calls `delete_db_instance` with the instance identifier and `SkipFinalSnapshot=True`. The script includes a waiter to ensure the instance is deleted and a final print statement. It also generates an EC2 client and calls `describe_security_groups` to get the security group ID, and then calls `delete_security_group` to delete the security group.

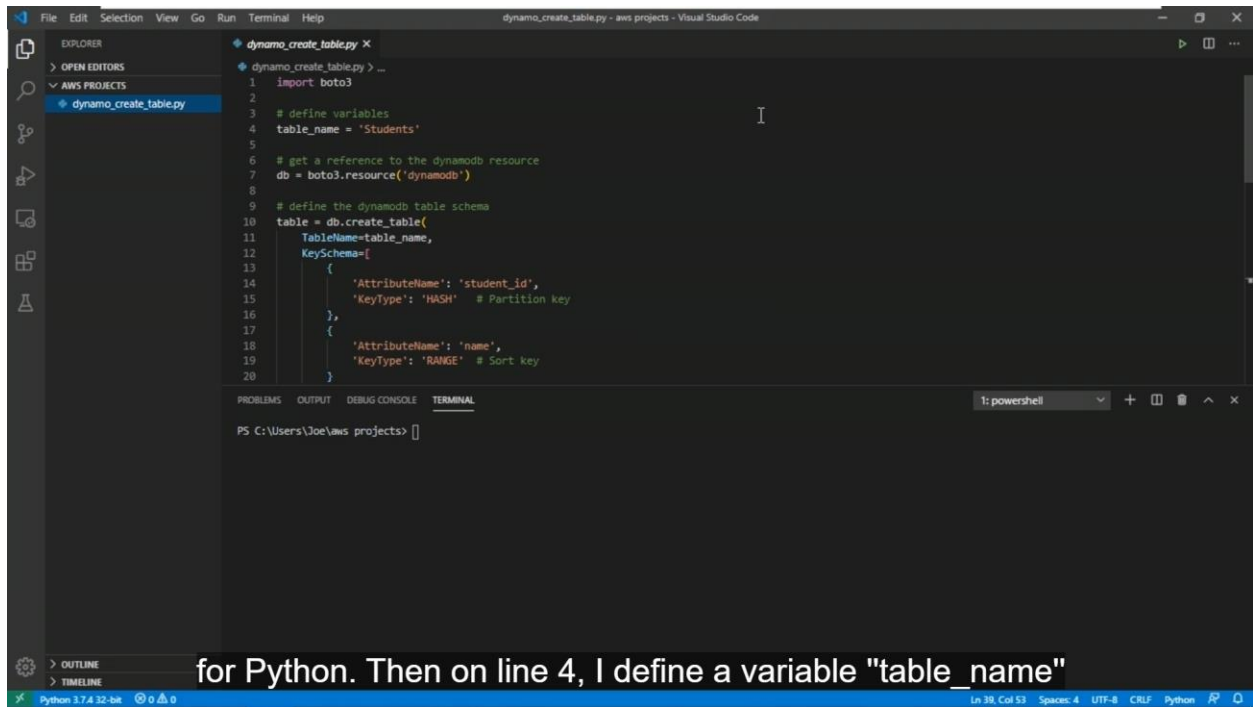
```
14 DBInstanceIdentifier=rds_identifier,
15 SkipFinalSnapshot=True)
16 print('Please wait while we terminate the "{}" RDS Instance'.format(rds_identifier))
17 waiter = cl.get_waiter('db_instance_deleted')
18 waiter.wait(DBInstanceIdentifier=rds_identifier,
19             WaiterConfig={'MaxAttempts': 50})
20 print('The "{}" RDS Instance has successfully been deleted!'.format(rds_identifier))
21
22 # generate EC2 client
23 cl = boto3.client('ec2')
24
25 # get the security group ID
26 res = cl.describe_security_groups(GroupNames=[sg_name])
27 sg_id = json.dumps(res['SecurityGroups'])[0]['GroupId'].replace("'", '')
28
29 # now we can delete the security group
30 res = cl.delete_security_group(GroupID=sg_id)
31 print('Security Group deleted!')
32
```

The terminal output shows the execution of the script:

```
PS C:\Users\joe\aws projects> python .\rds_del_db.py
Please wait while we terminate the "mysql-testdb" RDS Instance
The "mysql-testdb" RDS Instance has successfully been deleted!
Security Group deleted!
PS C:\Users\joe\aws projects>
```

OK. And now the message came back, "The 'mysql-testdb'

## Creating a Table in Amazon DynamoDB



```
File Edit Selection View Go Run Terminal Help
dynamo_create_table.py - aws projects - Visual Studio Code

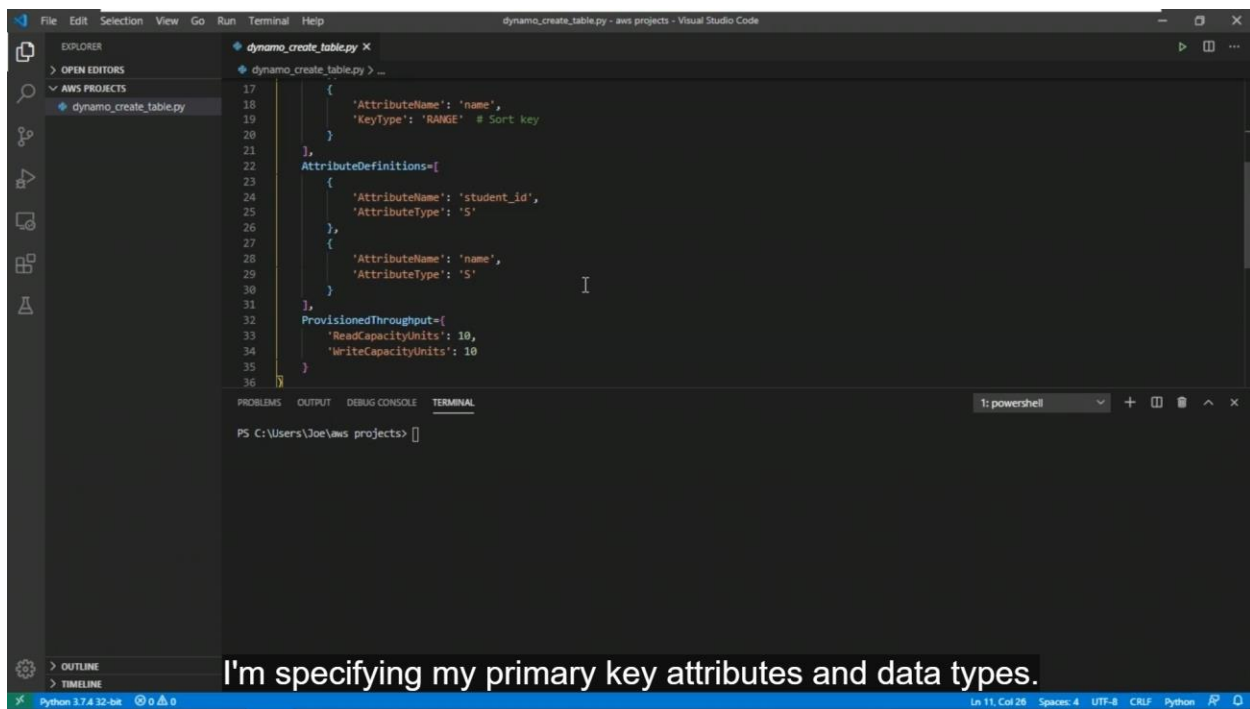
EXPLORER
> OPEN EDITORS
  AWS PROJECTS
    dynamo_create_table.py

dynamo_create_table.py x
dynamo_create_table.py > ...
1  import boto3
2
3  # define variables
4  table_name = 'Students'
5
6  # get a reference to the dynamodb resource
7  db = boto3.resource('dynamodb')
8
9  # define the dynamodb table schema
10 table = db.create_table(
11     TableName=table_name,
12     KeySchema=[
13         {
14             'AttributeName': 'student_id',
15             'KeyType': 'HASH' # Partition key
16         },
17         {
18             'AttributeName': 'name',
19             'KeyType': 'RANGE' # Sort key
20         }
21     ]
22 )

PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL
PS C:\Users\Joe\aws projects>
```

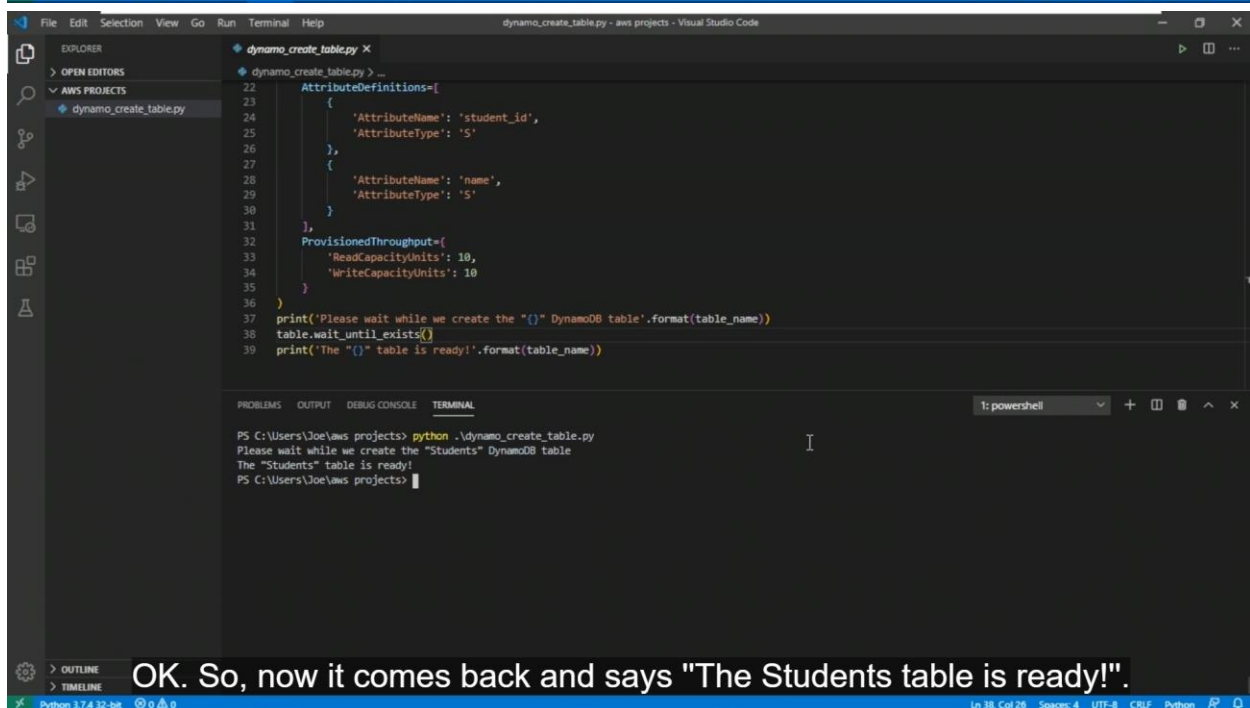
for Python. Then on line 4, I define a variable "table\_name"

Python 3.7.4 32-bit | Ln 39, Col 53 | Spaces: 4 | UTF-8 | CRLF | Python



```
17     {
18         'AttributeName': 'name',
19         'KeyType': 'RANGE' # Sort key
20     },
21 ],
22 AttributeDefinitions=[
23     {
24         'AttributeName': 'student_id',
25         'AttributeType': 'S'
26     },
27     {
28         'AttributeName': 'name',
29         'AttributeType': 'S'
30     }
31 ],
32 ProvisionedThroughput={
33     'ReadCapacityUnits': 10,
34     'WriteCapacityUnits': 10
35 }
36 ]
```

I'm specifying my primary key attributes and data types.



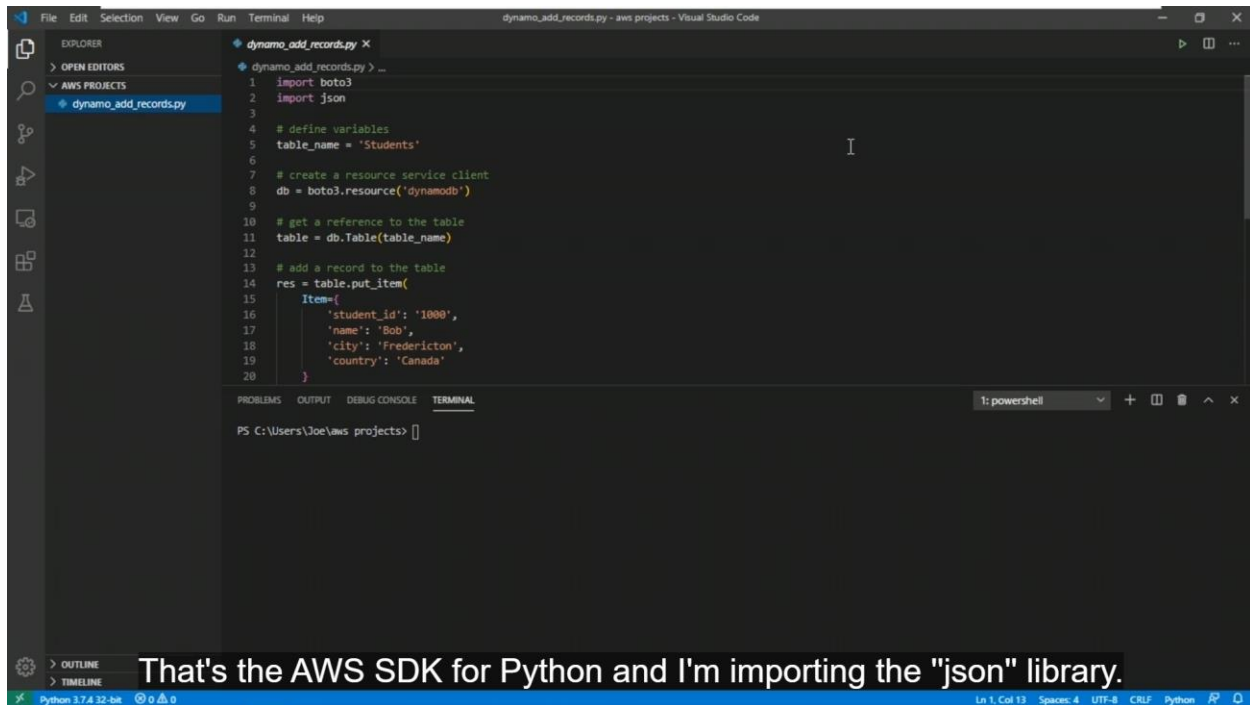
```
22 AttributeDefinitions=[
23     {
24         'AttributeName': 'student_id',
25         'AttributeType': 'S'
26     },
27     {
28         'AttributeName': 'name',
29         'AttributeType': 'S'
30     }
31 ],
32 ProvisionedThroughput={
33     'ReadCapacityUnits': 10,
34     'WriteCapacityUnits': 10
35 }
36 ]
37 print('Please wait while we create the "{}" DynamoDB table'.format(table_name))
38 table.wait_until_exists()
39 print('The "{}" table is ready!'.format(table_name))
```

PS C:\Users\Joe\aws projects> python dynamo\_create\_table.py  
Please wait while we create the "Students" DynamoDB table  
The "Students" table is ready!  
PS C:\Users\Joe\aws projects>

OK. So, now it comes back and says "The Students table is ready!"



## Adding Data to a DynamoDB Table



```
File Edit Selection View Go Run Terminal Help
dynamo_add_records.py - aws projects - Visual Studio Code

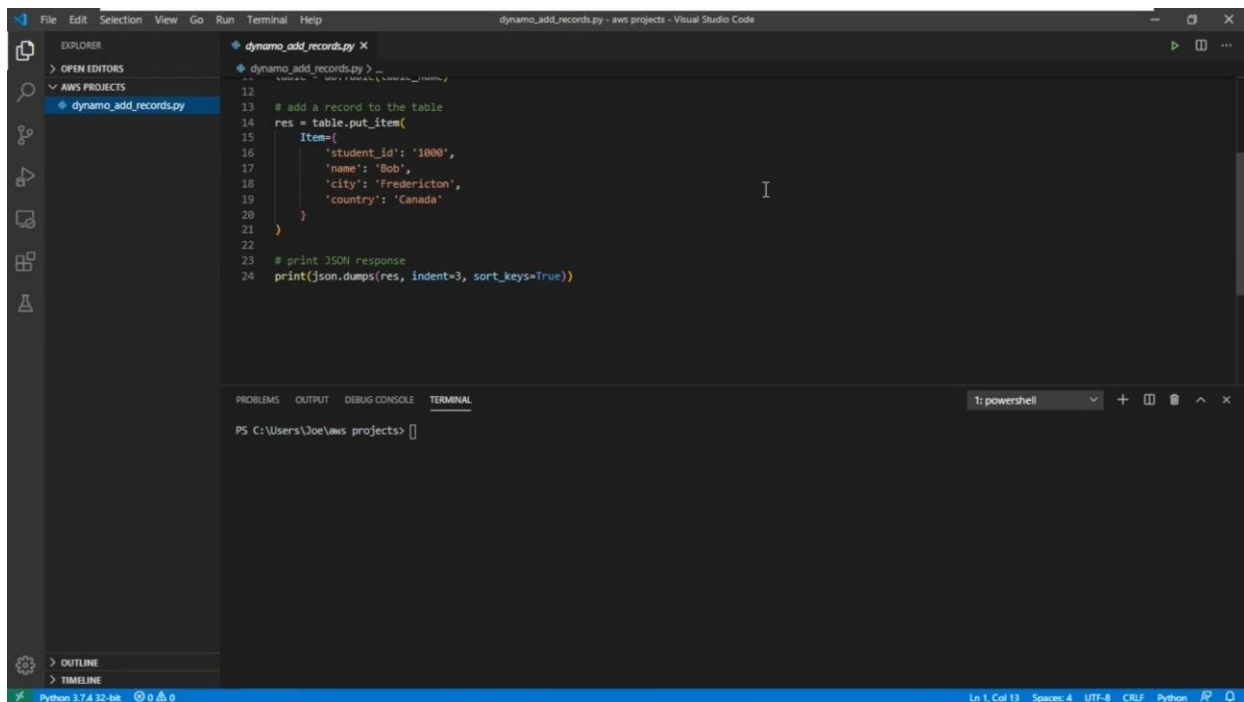
EXPLORER
  > OPEN EDITORS
  > AWS PROJECTS
    + dynamo_add_records.py

dynamo_add_records.py x
  dynamo_add_records.py > ...
  1 import boto3
  2 import json
  3
  4 # define variables
  5 table_name = 'Students'
  6
  7 # create a resource service client
  8 db = boto3.resource('dynamodb')
  9
 10 # get a reference to the table
 11 table = db.Table(table_name)
 12
 13 # add a record to the table
 14 res = table.put_item(
 15     Item={
 16         'student_id': '1000',
 17         'name': 'Bob',
 18         'city': 'Fredericton',
 19         'country': 'Canada'
 20     }
 21 )

PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL
PS C:\Users\Joe\aws projects>
```

That's the AWS SDK for Python and I'm importing the "json" library.

Python 3.7.4 32-bit 0 0 0 Ln 1, Col 13 Spaces: 4 UTF-8 CRLF Python

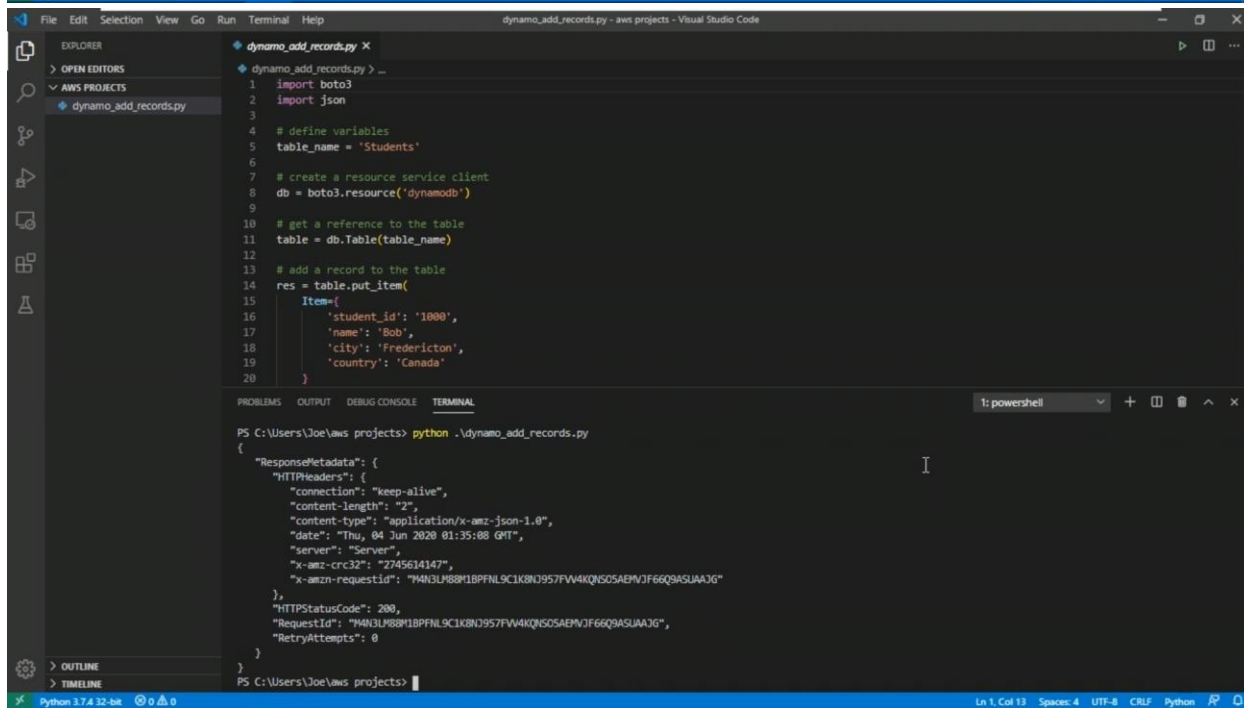


```
File Edit Selection View Go Run Terminal Help
dynamo_add_records.py - aws projects - Visual Studio Code

EXPLORER
> OPEN EDITORS
AWS PROJECTS
  dynamo_add_records.py

dynamo_add_records.py
12
13 # add a record to the table
14 res = table.put_item(
15     Item={
16         'student_id': '1000',
17         'name': 'Bob',
18         'city': 'Fredericton',
19         'country': 'Canada'
20     }
21 )
22
23 # print JSON response
24 print(json.dumps(res, indent=3, sort_keys=True))

PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL
1: powershell
PS C:\Users\Joe\aws projects>
```



```
File Edit Selection View Go Run Terminal Help
dynamo_add_records.py - aws projects - Visual Studio Code

EXPLORER
> OPEN EDITORS
AWS PROJECTS
  dynamo_add_records.py

dynamo_add_records.py
1 import boto3
2 import json
3
4 # define variables
5 table_name = 'Students'
6
7 # create a resource service client
8 db = boto3.resource('dynamodb')
9
10 # get a reference to the table
11 table = db.Table(table_name)
12
13 # add a record to the table
14 res = table.put_item(
15     Item={
16         'student_id': '1000',
17         'name': 'Bob',
18         'city': 'Fredericton',
19         'country': 'Canada'
20     }
21 )

PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL
1: powershell
PS C:\Users\Joe\aws projects> python .\dynamo_add_records.py
{
  "ResponseMetadata": {
    "HTTPHeaders": {
      "connection": "keep-alive",
      "content-length": "2",
      "content-type": "application/x-amz-json-1.0",
      "date": "Thu, 04 Jun 2020 01:35:00 GMT",
      "server": "Server",
      "x-amz-crc32": "2745614147",
      "x-amzn-requestid": "f4N3U1B8B1BPFNL9C1KBN957FW4KQNS0SAEPVJF66Q9ASUAAJG"
    },
    "HTTPStatusCode": 200,
    "RequestId": "f4N3U1B8B1BPFNL9C1KBN957FW4KQNS0SAEPVJF66Q9ASUAAJG",
    "RetryAttempts": 0
  }
}
PS C:\Users\Joe\aws projects>
```

# Querying a DynamoDB Table

The image consists of two screenshots of a Visual Studio Code editor window, showing the process of querying a DynamoDB table using Python and boto3.

**Top Screenshot:** The editor shows a file named `dynamo_query_table.py`. The code defines a table named 'Students' and performs a query for a student with ID '1000'. The response is stored in `res`, but the output is currently empty.

```
1 import boto3
2 from boto3.dynamodb.conditions import Key
3 import json
4
5 # define variables
6 table_name = 'Students'
7
8 # creating a resource service client
9 db = boto3.resource('dynamodb')
10
11 # get a reference to the table
12 table = db.Table(table_name)
13
14 # perform query
15 res = table.query(KeyConditionExpression=Key('student_id').eq('1000'))
16
17 # print the response
18 if len(res['Items']) == 0:
19     print('Nothing found!')
20 else:
```

The terminal shows the command prompt: `PS C:\Users\Joe\aws projects>`.

**Bottom Screenshot:** The code is updated to print the response in a formatted JSON structure. The terminal output shows the result of the query.

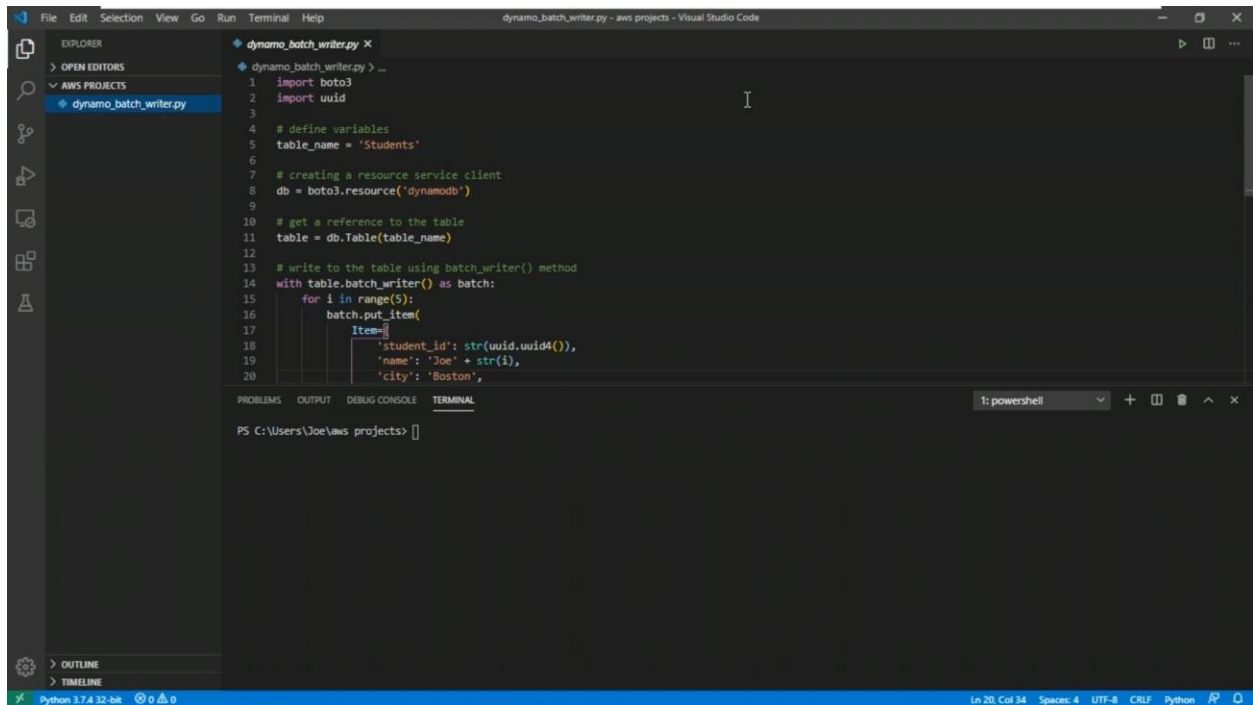
```
21     print(json.dumps(res['Items'], indent=3, sort_keys=True))
```

The terminal shows the command `python .\dynamo_query_table.py` being executed, resulting in the following output:

```
[
  {
    "city": "Fredericton",
    "country": "Canada",
    "name": "Bob",
    "student_id": "1000"
  }
]
```

The status bar at the bottom indicates the file is using Python 3.7.4 32-bit, with 15 lines and 16 columns selected, in UTF-8 encoding with CRLF line endings.

## Using a Batch Writer against a DynamoDB Table



```
File Edit Selection View Go Run Terminal Help
dynamo_batch_writer.py - aws projects - Visual Studio Code

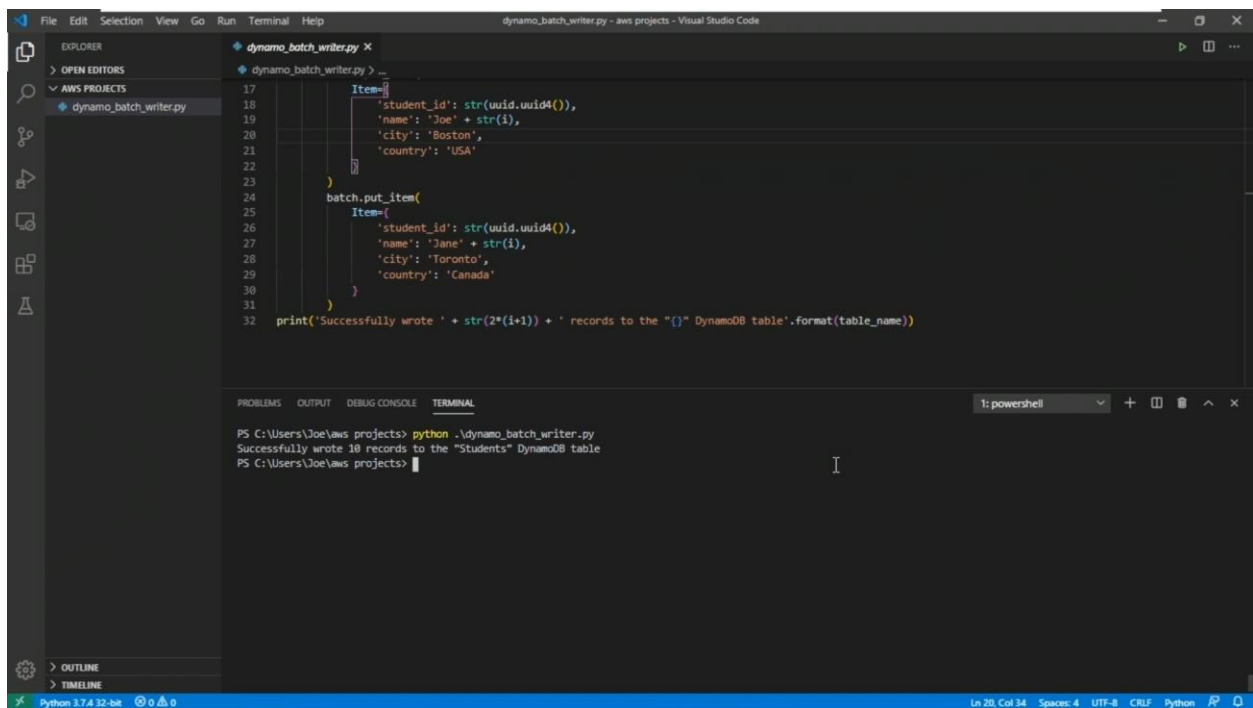
EXPLORER
  OPEN EDITORS
  AWS PROJECTS
    dynamo_batch_writer.py

dynamo_batch_writer.py
1 import boto3
2 import uuid
3
4 # define variables
5 table_name = 'Students'
6
7 # creating a resource service client
8 db = boto3.resource('dynamodb')
9
10 # get a reference to the table
11 table = db.Table(table_name)
12
13 # write to the table using batch_writer() method
14 with table.batch_writer() as batch:
15     for i in range(5):
16         batch.put_item(
17             Item={
18                 'student_id': str(uuid.uuid4()),
19                 'name': 'Joe' + str(i),
20                 'city': 'Boston',
                
```

PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL

PS C:\Users\Joe\aws projects>

Python 3.7.4 32-bit



```
File Edit Selection View Go Run Terminal Help
dynamo_batch_writer.py - aws projects - Visual Studio Code

EXPLORER
  OPEN EDITORS
  AWS PROJECTS
    dynamo_batch_writer.py

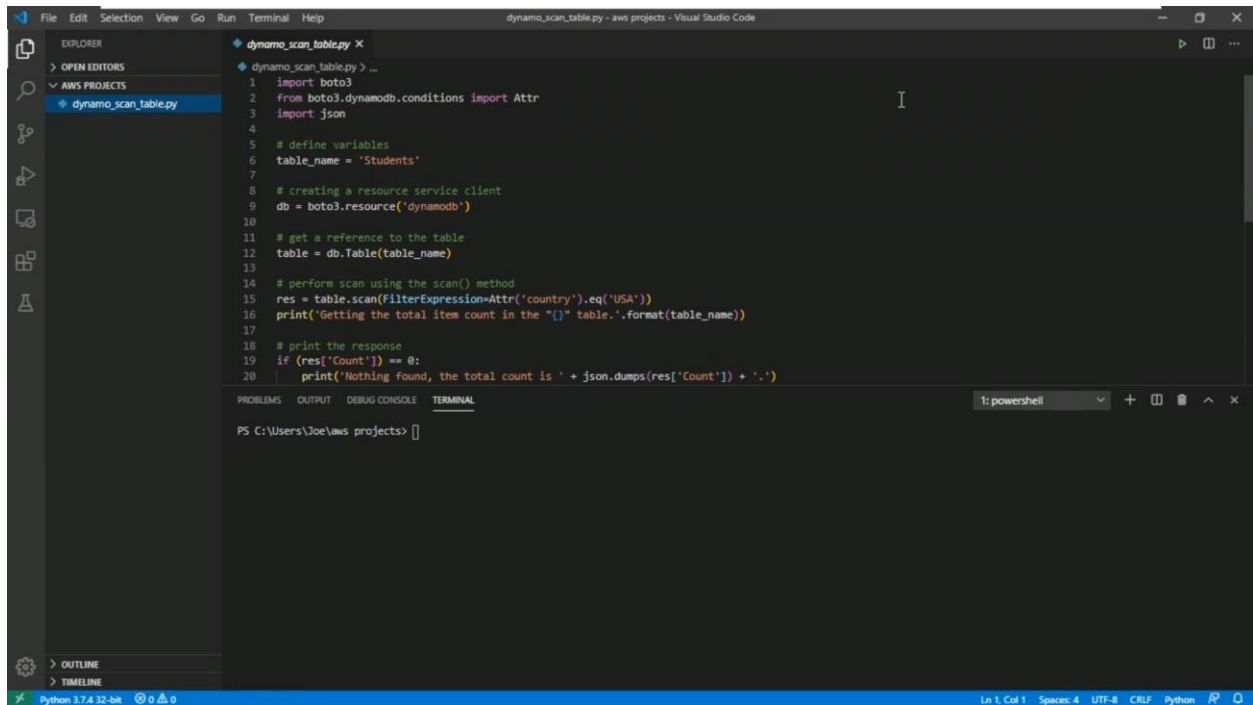
dynamo_batch_writer.py
17
18         'student_id': str(uuid.uuid4()),
19         'name': 'Joe' + str(i),
20         'city': 'Boston',
21         'country': 'USA'
22     }
23     batch.put_item(
24         Item={
25             'student_id': str(uuid.uuid4()),
26             'name': 'Jane' + str(i),
27             'city': 'Toronto',
28             'country': 'Canada'
29         }
30     )
31
32 print('Successfully wrote ' + str(2*(i+1)) + ' records to the "{}" DynamoDB table'.format(table_name))

PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL

PS C:\Users\Joe\aws projects> python .\dynamo_batch_writer.py
Successfully wrote 10 records to the "Students" DynamoDB table
PS C:\Users\Joe\aws projects>
```

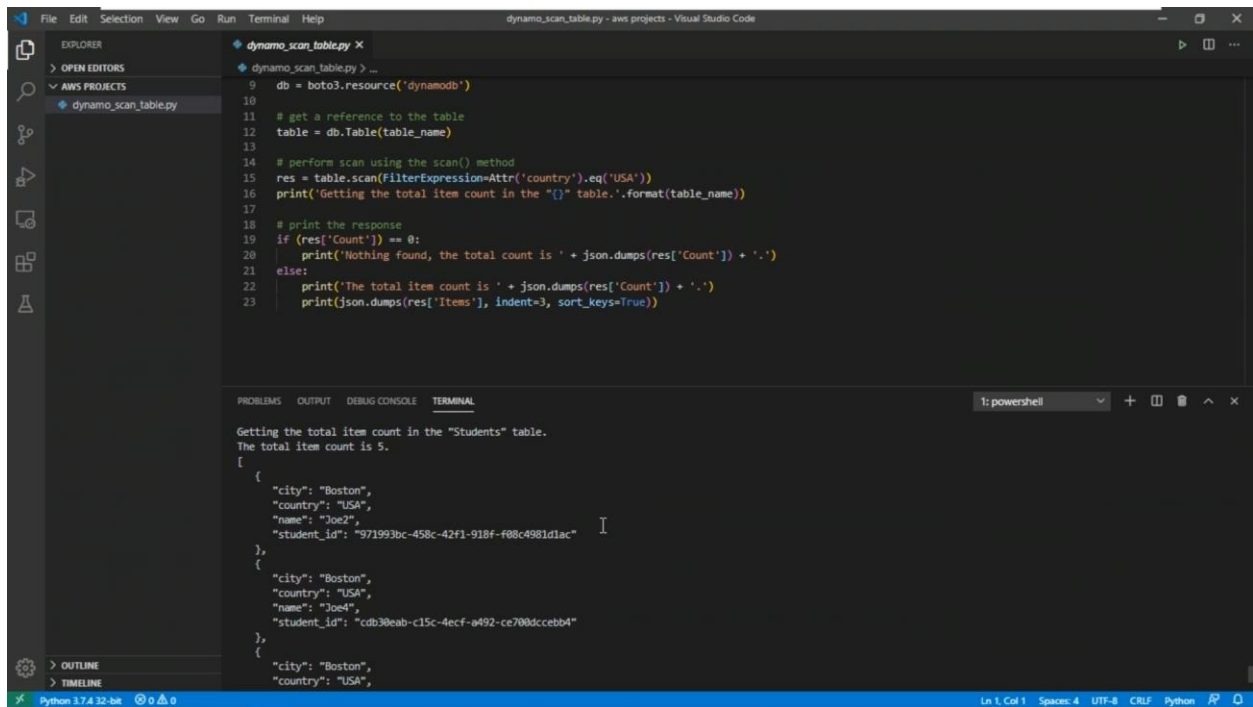
Python 3.7.4 32-bit

## Scanning a DynamoDB Table



The screenshot shows the Visual Studio Code editor with a file named `dynamo_scan_table.py` open. The code is written in Python and uses the `boto3` library to interact with AWS DynamoDB. The code defines a table name, creates a resource service client, gets a reference to the table, and performs a scan using the `scan()` method. The scan is filtered by the `country` attribute being equal to `'USA'`. The code prints the total item count in the table and the response. The terminal window at the bottom shows the command prompt `PS C:\Users\Joe\aws projects>`.

```
1 import boto3
2 from boto3.dynamodb.conditions import Attr
3 import json
4
5 # define variables
6 table_name = 'Students'
7
8 # creating a resource service client
9 db = boto3.resource('dynamodb')
10
11 # get a reference to the table
12 table = db.Table(table_name)
13
14 # perform scan using the scan() method
15 res = table.scan(FilterExpression=Attr('country').eq('USA'))
16 print('Getting the total item count in the "{}" table.'.format(table_name))
17
18 # print the response
19 if (res['Count']) == 0:
20     print('Nothing found, the total count is ' + json.dumps(res['Count']) + '.')
21
```

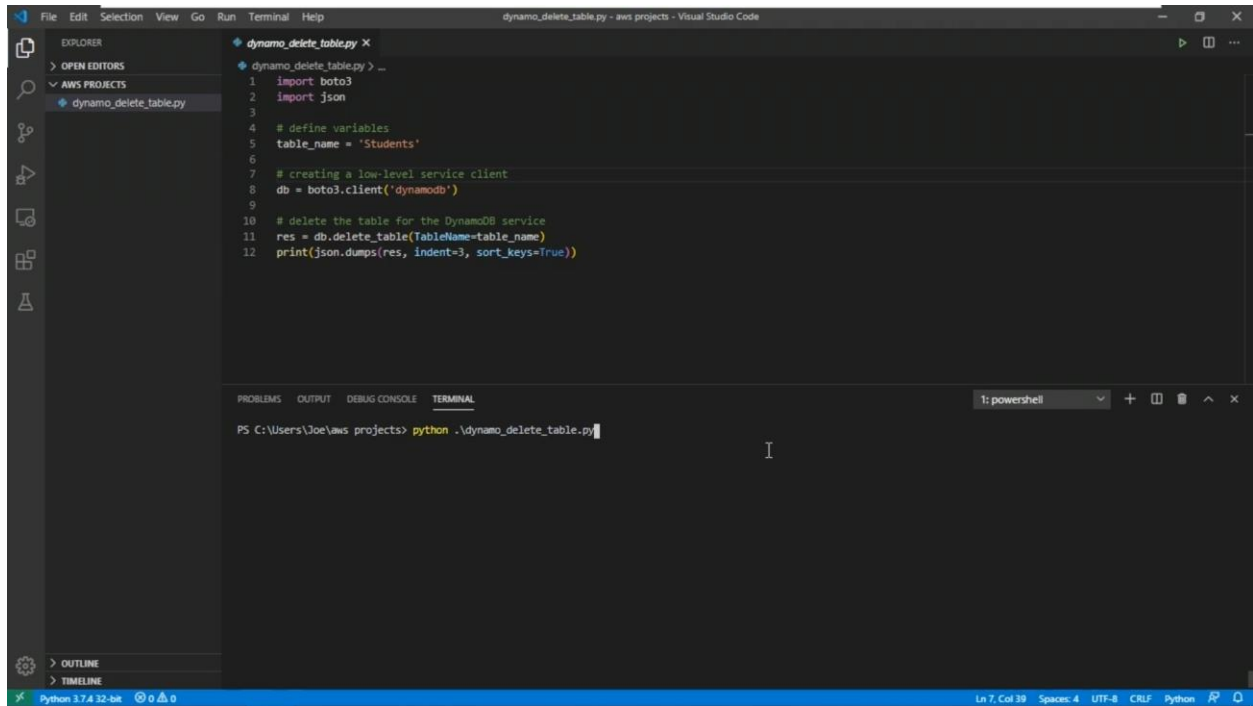


The screenshot shows the Visual Studio Code editor with the same file `dynamo_scan_table.py` open. The code is identical to the previous screenshot, but the terminal window at the bottom now shows the output of the scan operation. The output indicates that the total item count in the `Students` table is 5, and it displays a list of items with their attributes.

```
Getting the total item count in the "Students" table.
The total item count is 5.
[
  {
    "city": "Boston",
    "country": "USA",
    "name": "Joe2",
    "student_id": "971993bc-458c-42f1-918f-f08c4981d1ac"
  },
  {
    "city": "Boston",
    "country": "USA",
    "name": "Joe4",
    "student_id": "c3b30eab-c15c-4ecf-a492-ce700dcebb4"
  },
  {
    "city": "Boston",
    "country": "USA",

```

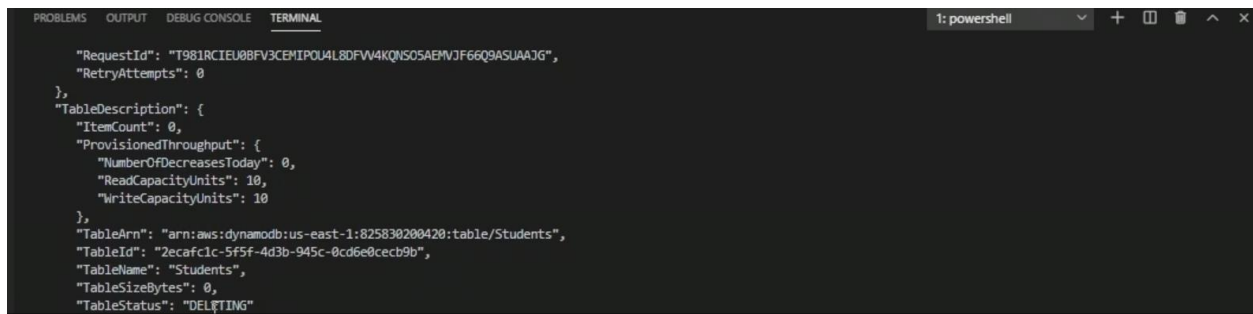
## Deleting a DynamoDB Table



The screenshot shows the Visual Studio Code interface with a file explorer on the left and a code editor in the center. The file explorer shows a project named 'AWS PROJECTS' with a file named 'dynamo\_delete\_table.py'. The code editor displays the following Python script:

```
1 import boto3
2 import json
3
4 # define variables
5 table_name = 'Students'
6
7 # creating a low-level service client
8 db = boto3.client('dynamodb')
9
10 # delete the table for the DynamoDB service
11 res = db.delete_table(TableName=table_name)
12 print(json.dumps(res, indent=3, sort_keys=True))
```

Below the code editor is a terminal window with the command prompt 'PS C:\Users\Joe\aws projects> python .\dynamo\_delete\_table.py'.



The screenshot shows the terminal output of the command executed in the previous screenshot. The output is a JSON object representing the response from the delete\_table command:

```
{
  "RequestId": "T981RCIEU08FV3CEIPOU4L8DFW4KQNS05AEMVJF66Q9ASUAAJG",
  "RetryAttempts": 0,
  "TableDescription": {
    "ItemCount": 0,
    "ProvisionedThroughput": {
      "NumberOfDecreasesToday": 0,
      "ReadCapacityUnits": 10,
      "WriteCapacityUnits": 10
    },
    "TableArn": "arn:aws:dynamodb:us-east-1:825830200420:table/Students",
    "TableId": "2eacfc1c-5f5f-4d3b-945c-0cd6e0cecb9b",
    "TableName": "Students",
    "TableSizeBytes": 0,
    "TableStatus": "DELETING"
  }
}
```

1. Which relational database engines are available through Amazon RDS?
  - **MySQL**
  - **PostgreSQL**
2. When deleting an RDS instance, which waiter should be used in Python code to wait for the database instance to be identified as deleted by Amazon RDS?
  - **db\_instance\_deleted**
3. What tool does AWS Data Migration Service (DMS) use for heterogeneous migrations to ensure compatible schema in the target database?
  - **Schema Conversion Tool (SCT)**
4. You've installed boto3 and PyMySQL using Python's pip package manager. What single command can be used to verify details about the versions of these packages installed?
  - **pip show boto3 pymysql**
5. While using the create\_table() method to create a table in Amazon DynamoDB, what information must be specified in the call to create\_table() for a provisioned table?
  - **Table name**
  - **Primary key data types**
  - **Throughput settings**
  - **Primary key attributes**
6. Complete the code snippet to query a table in Amazon RDS.  
# query the table within a try block  
try:  
    sql = "SELECT \* FROM `Students` "  
    cursor.execute(sql)  
    res = cursor.<missing code>  
    print('Querying table...')  
    print(res)  
except MySQLdb.Error as e:  
    print('Error: {}'.format(e))

```

        print('Something went wrong!')
    finally:
        db.close()
-   fetchall()

```

7. Complete the code snippet to generate a cursor that will be used to create a table in an Amazon RDS MySQL database.

```
# connect to the target database
```

```
db = mysqlldb.connect(
    host=rds_endpoint, user=user_name,
    password=user_password, database=db_name)
```

```
cursor = db.<missing code>()
```

```
# create the table within a try block
```

```
try:
```

```
    cursor.execute(
        "CREATE TABLE Students " +
        "(st_id INT NOT NULL AUTO_INCREMENT, name VARCHAR(20) NOT NULL," +
        "PRIMARY KEY (`student_id` ))")
    print('Table successfully created!')
```

```
except mysqlldb.Error as e:
```

```
    print('Error: {}'.format(e))
```

```
finally:
```

```
    db.close()
-   cursor
```

8. Which nonrelational databases are available from AWS?

- **Amazon DocumentDB**
- **Amazon DynamoDB**



9. Complete the code snippet to generate the inbound rules for a Security Group on AWS Cloud using the AWS SDK for Python.

```
# add rules for the security group
res = cl.<missing code>(
    CidrIp=ip_cidr,
    FromPort=3306,
    GroupName=sg_name,
    ToPort=3306,
    IpProtocol='tcp')
print("Security Group successfully created!")
```

- **authorize\_security\_group\_ingress**

10. Which statement is used to import the Key class so that a query condition can be related to the key of an item.

- **from boto3.dynamodb.conditions import Key**

11. Which types of NoSQL databases are available from AWS?

- **Columnar**
- **Documents**

12. Which statement is used to import the appropriate class so that we can add conditions to a scanning operation against a DynamoDB table?

- **from boto3.dynamodb.conditions import Attr**

13. What are the two main strategies of caching?

- **Write through**
- **Lazy loading**

14. Complete the code snippet to create a resource service client that can be used on the DynamoDB service to add items to a table.

```
import boto3
import json
```

```
# define variables
table_name = 'Students'
```

```
# create a resource service client
db = boto3.resource('dynamodb')
```

- **resource('dynamodb')**

15. Complete the code snippet to delete a DynamoDB table using the delete\_table() method.

```
import boto3
import json
```

```
# define variables
table_name = 'Students'
```

```
# creating a low-level service client
db = <missing code>
```

```
# delete the table for the DynamoDB service
res = db.delete_table(TableName=table_name)
print(json.dumps(res, indent=3, sort_keys=True))
```

- **boto3.client('dynamodb')**

16. Which characteristics accurately describe Amazon DynamoDB?

- **Features automatic scaling**
- **Fully managed database service**

17. What is the only DB instance size available if you want to use RDS free tier?

- **db.t2.micro**

18. Complete the code snippet to generate items in a DynamoDB table with the batch\_writer() method.

```
# write to the table using batch_writer() method
```

```
with table.batch_writer() as batch:
```

```
    for i in range(5):
```

```
        batch.<missing code>(  
            Item={
```

```
                'student_id': str(uuid.uuid4()),
```

```
                'name': 'Joe' + str(i),
```

```
                'city': 'Boston',
```

```
                'country': 'USA'
```

```
            }  
        )
```

- **put\_item**

19. You've installed MySQL Workbench 8.0 CE and you are configuring a new connection. In which field on the Setup New Connection dialog do you specify the Amazon RDS instance endpoint?

- **Hostname**

20. Which statements accurately describe relational database characteristics?

- **Relational databases incorporate integrity constraints**
- **Relational databases use the concept of transactions**